

CSC 374/407: Computer Systems II

Lecture 2

Joseph Phillips
De Paul University

2014 January 3

Copyright © 2011 Joseph Phillips
All rights reserved

Reading

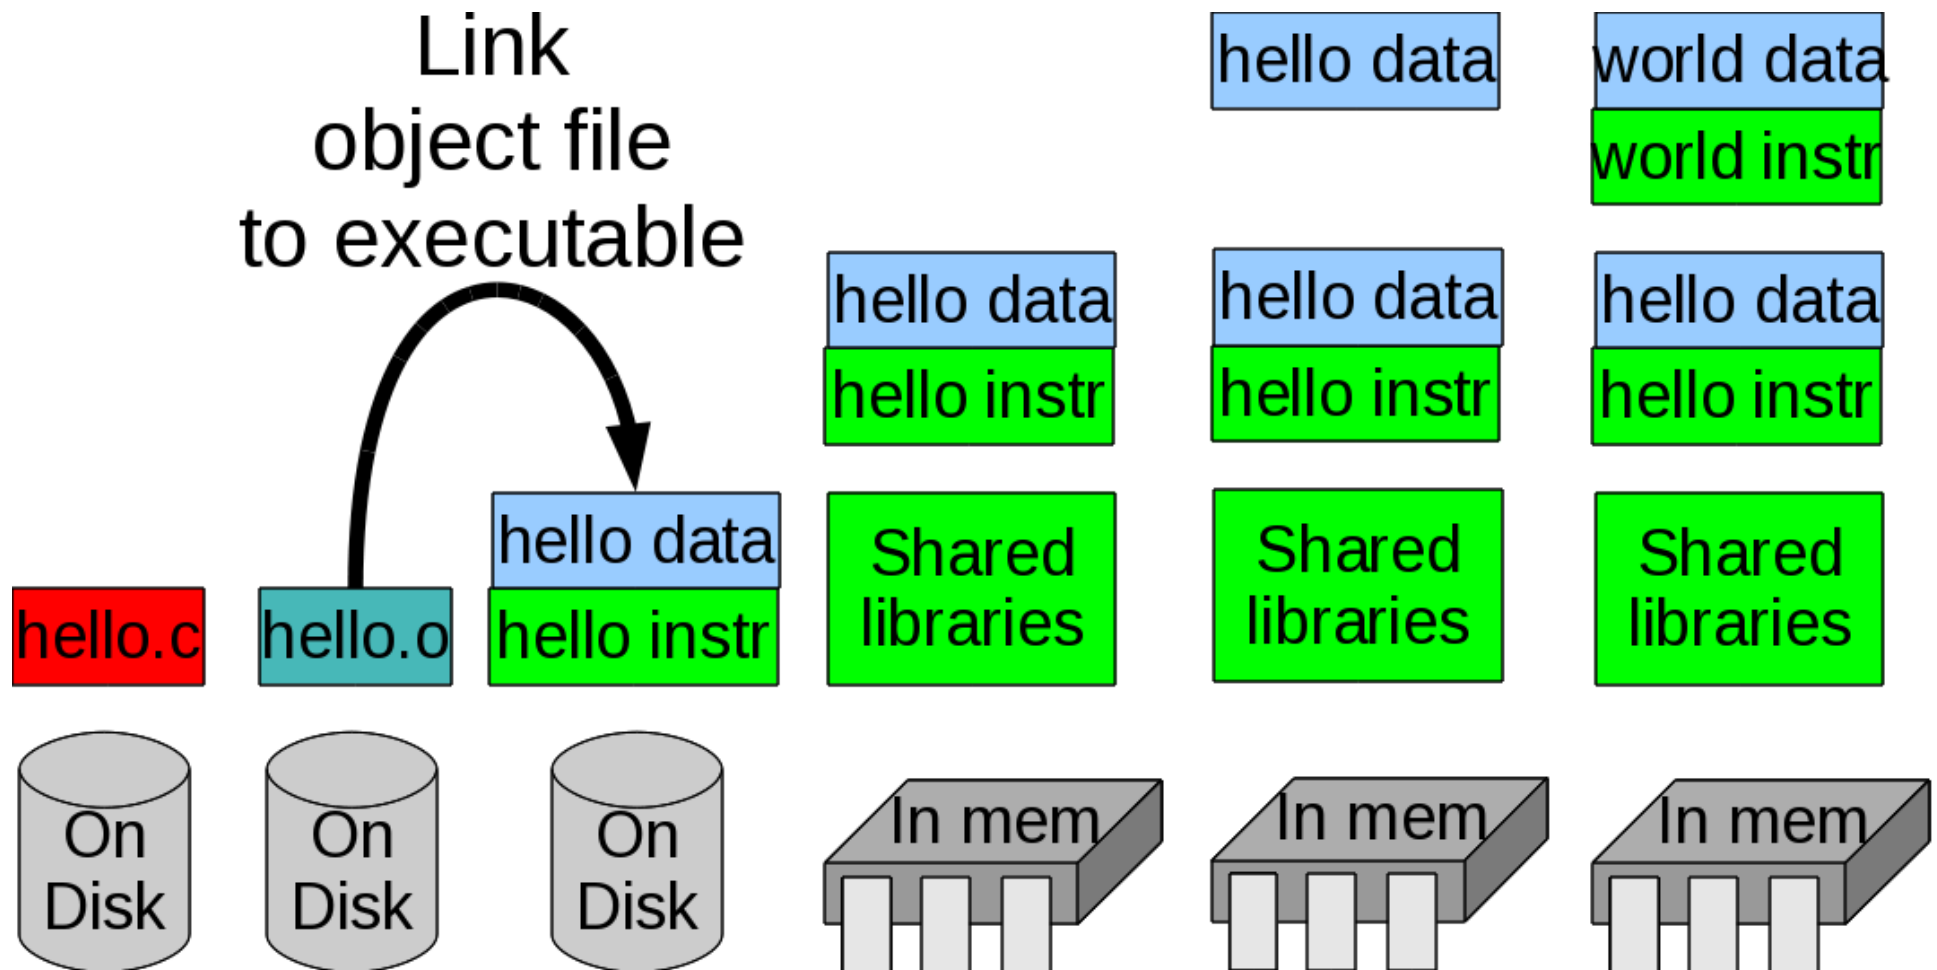
- ♦ Bryant & O'Hallaron “*Computer Systems, 2nd Ed.*”
 - ♦ Chapter 7: Linking
- ♦ Hoover “*System Programming*”
 - ♦ Program Management 6
 - ♦ Libraries 8.1-8.3, 8.6, 8.7

Topics

- ◆ Linkers
- ◆ ELF
- ◆ Dynamic Linking

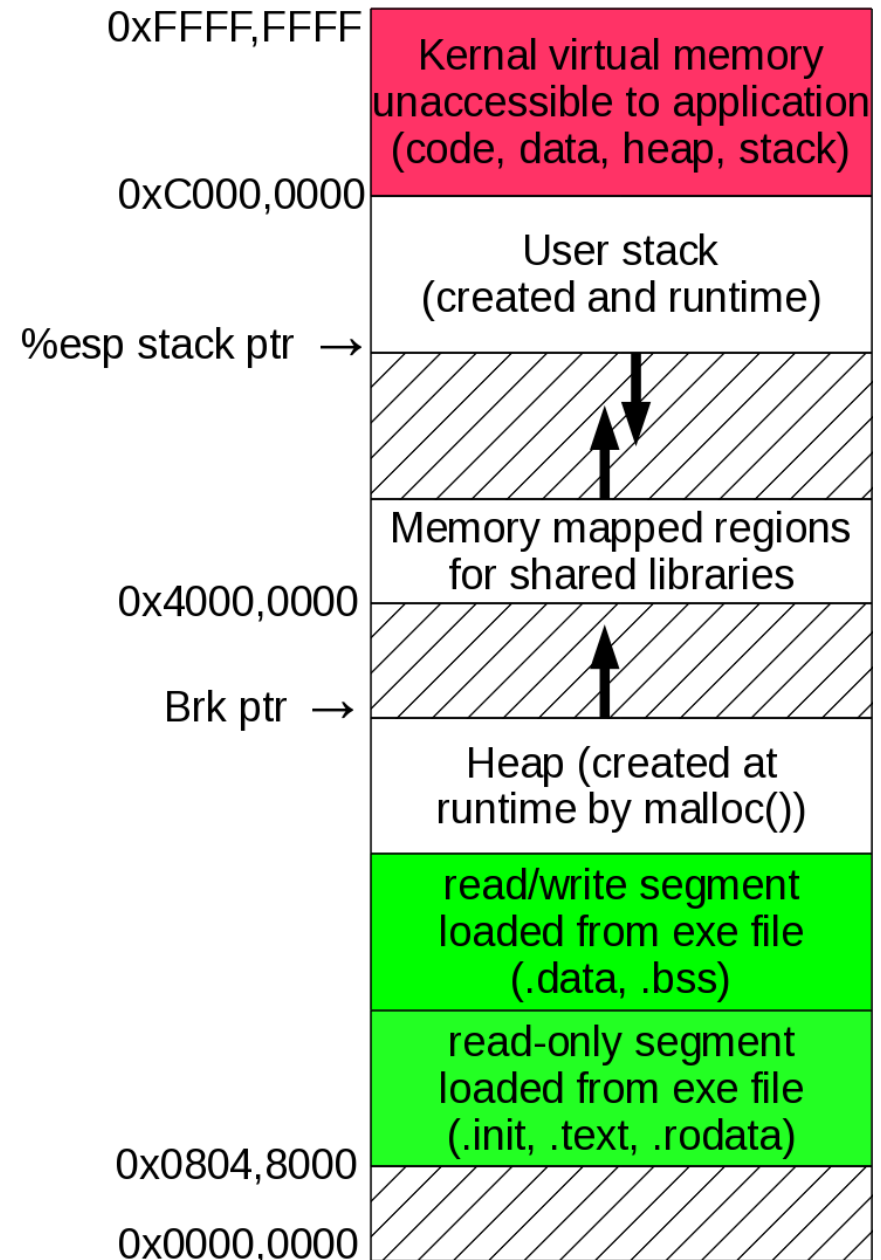
Today's topic (in time)

Linking object files into executable files



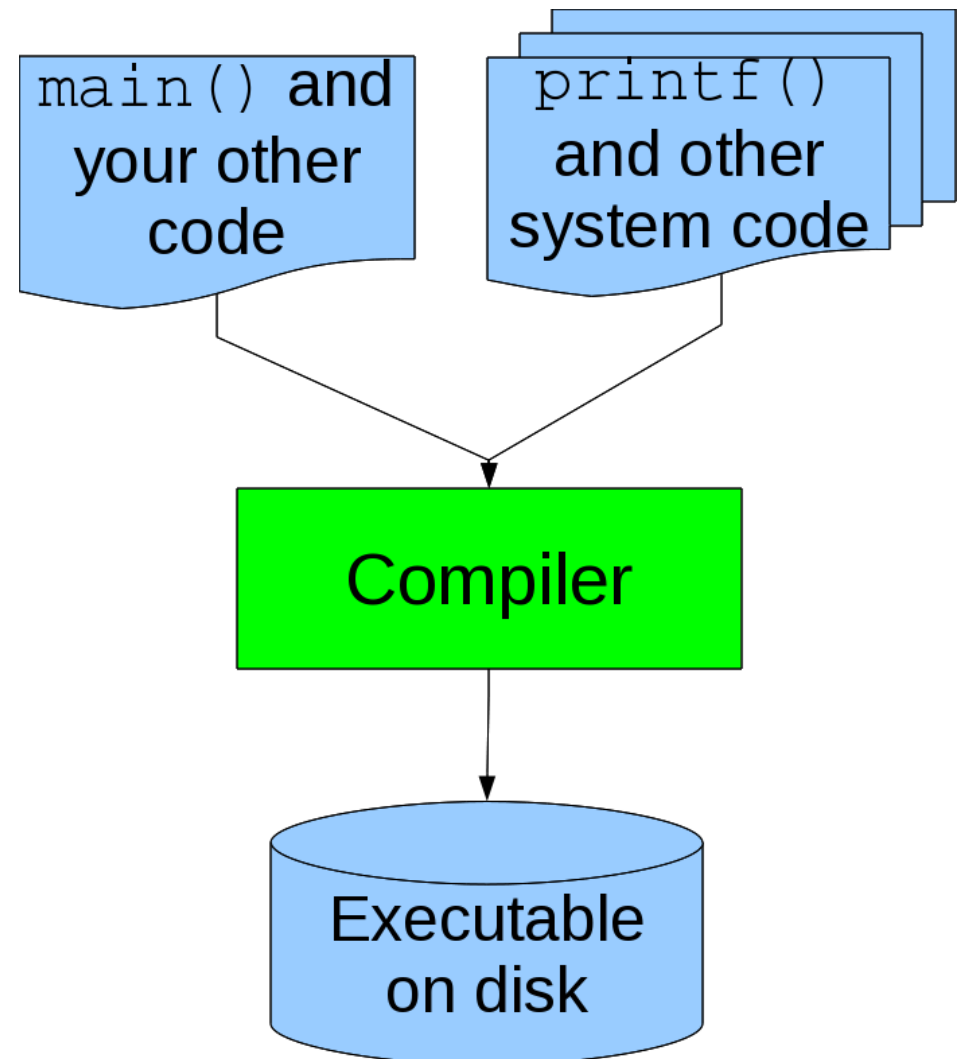
Today's topic (in space)

- Parts of the program the compiler actually has to create



How would *you* write a compiler?

- ♦ The direct approach:
- ♦ One program to compile ***everything***:
 - ♦ ***Everything*** includes your code (naturally!)
 - ♦ ***Everything*** includes standard library code too (*hmm?*)

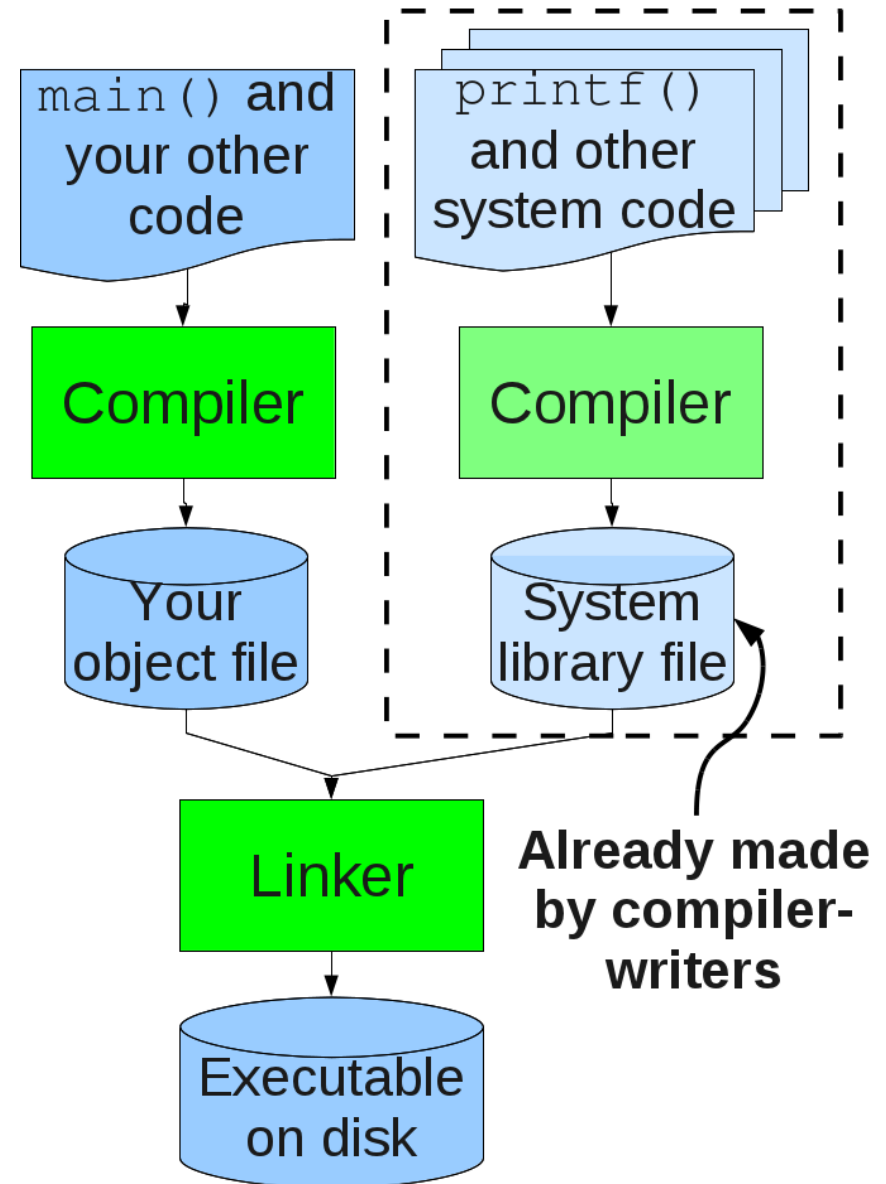


What?!? Re-compiling library code again and again?!?

- ◆ Sounds like a bad idea because
 1. *It's a waste of time*: mostly it doesn't change.
 2. *So much for hiding implementation*: When given easy-access to the source code some folks just can't stop themselves from trying to “optimize” their own code after seeing how the standard library code actually works. *Even worse* they can start to muck with it . . .
- ◆ So, let's give them the library code in a *post-compiled* but *pre-integrated* into a program . . .

Enter the linker!

- Compile standard code into libraries.
- Distribute libraries to user-programmers
- Programmers compile their own code to object files
- ***Linkers combine object files + libraries into executable files***



Linkers give you . . .

♦ **Modularity**

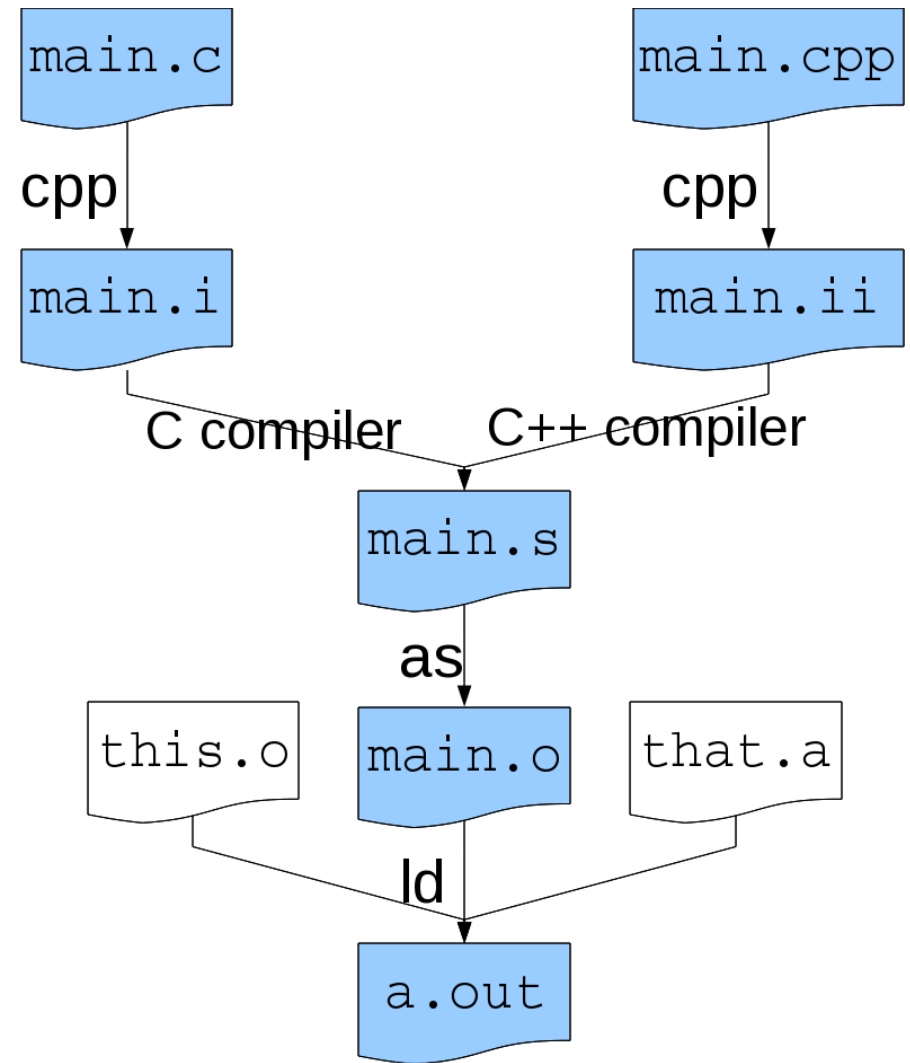
- ♦ You can write your program as small files can compile them individually
- ♦ You can use the libraries of others or build your own

♦ **Efficiency**

- ♦ ***Time***: No need to re-compile everything if change one small file
- ♦ ***Space***: Can incorporate from libraries only code that's actually used

There's a linker under the hood?

- **You** “*But all I have to say is `gcc main.c`*”
- **Prof Joe** “*Yes, 'cuz it does a number of things behind your back:*”
 - `cpp`: the C preprocessor
 - The C/C++ compiler proper
 - `as`: the assembler
 - `ld`: the linker



Let's try!

- ◆ You can tell `gcc` and `g++` to stop along the way:
 - ◆ `gcc -E test.c` # stop after preprocessing
 - ◆ Output is post-processed C code to `stdout`
 - ◆ `gcc -S test.c` # stop after compiling
 - ◆ Output is assembly code in `test.s`
 - ◆ `gcc -c test.c` # stop after assembling
 - ◆ Output is object file `test.o`
 - ◆ `gcc test.c` # Go ***all the way*** baby!
 - ◆ Output is executable file `a.out`

How do Linkers Spend Their Time?

To merge object files they must:

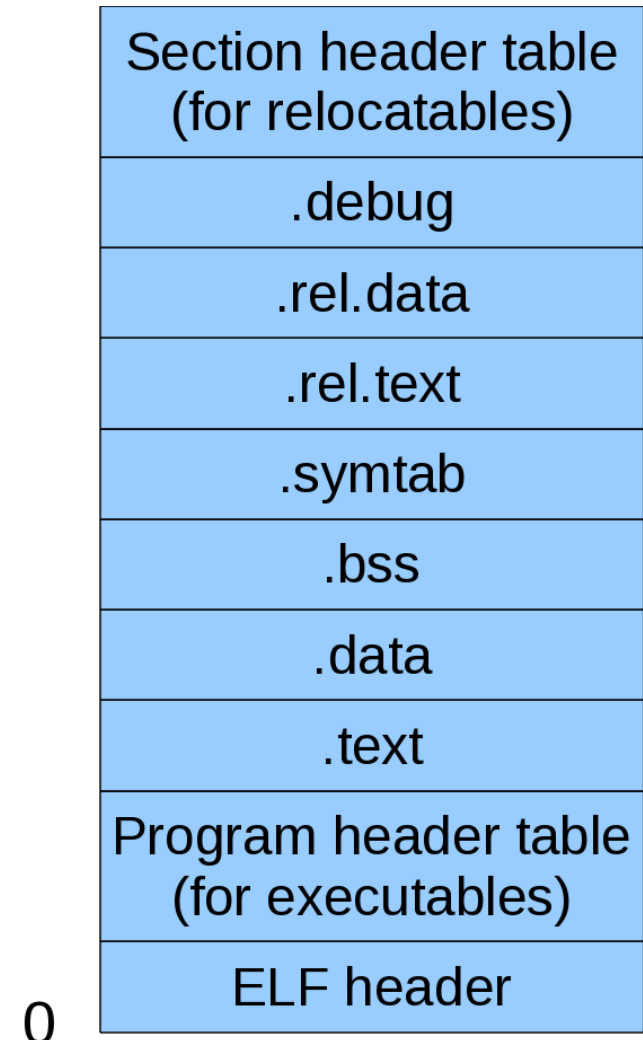
- ◆ Resolve “external references:”

```
extern int externallyDefinedVar;  
extern int externallyDefinedFunc(int);  
int      locallyDefinedGlobalVar;  
int main ()  
{  
    . . .  
    externallyDefinedVar=externallyDefinedFunc(5);  
}
```

- ◆ Move all symbols from several object files to distinct places in executable
- ◆ Update symbol locations
- ◆ “Symbol” means either named code or data

Executable and Link Format (ELF)

- ◆ Assembler, Linker and Program Loader all need one format to describe object files and executables
- ◆ ELF derived from AT&T Unix System V and BSD
- ◆ Unifies description of:
 - ◆ Relocatable object files (.o)
 - ◆ Executables
 - ◆ Shared object files (.so)



What's in an ELF?

- ◆ ELF header:
 - ◆ “Magic number” (Am I for Intel 386? Sparc Solaris? Something else?)
 - ◆ Important for shared disk systems
 - ◆ Byte ordering
 - ◆ Big or little endian?
 - ◆ Type
 - ◆ .o? Executable? .so?

Section header table (for relocatables)
.debug
.rel.data
.rel.text
.symtab
.bss
.data
.rodata
.text
Program header table (for executables)
ELF header

0

What's in an ELF?

- ◆ Program header:
 - ◆ OS-specific info like
 - ◆ Page size
 - ◆ Virtual address mem segments
 - ◆ Segment sizes

0	Section header table (for relocatables)
	.debug
	.rel.data
	.rel.text
	.symtab
	.bss
	.data
	.rodata
	.text
	Program header table (for executables)
	ELF header

What's in an ELF?

- ♦ `.text`
 - ♦ Code (functions)
- ♦ `.rodata`
 - ♦ Read-Only consts (e.g. strings)
- ♦ `.data`
 - ♦ Global and static vars initialized to other than 0
- ♦ `.bss`
 - ♦ Global and static vars that start out 0
 - ♦ (Why not in `.data`?)

0	Section header table (for relocatables)
	<code>.debug</code>
	<code>.rel.data</code>
	<code>.rel.text</code>
	<code>.symtab</code>
	<code>.bss</code>
	<code>.data</code>
	<code>.rodata</code>
	<code>.text</code>
	Program header table (for executables)
	ELF header

What's in an ELF?

- ◆ `.symtab`
 - ◆ Symbol table holds names and locations of functions, variables and sections
- ◆ `.rel.text`
 - ◆ Relocatable info for `.text` (which `.text` addresses need to be changed and how)
- ◆ `.rel.data`
 - ◆ Ditto for `.data`

Section header table (for relocatables)	
	<code>.debug</code>
	<code>.rel.data</code>
	<code>.rel.text</code>
	<code>.symtab</code>
	<code>.bss</code>
	<code>.data</code>
	<code>.rodata</code>
	<code>.text</code>
Program header table (for executables)	
	ELF header

0

What's in an ELF?

- ♦ `.debug`
 - ♦ Information for `gdb` debugger
 - ♦ Generated when compile with `-g` option

Section header table (for relocatables)	
	<code>.debug</code>
	<code>.rel.data</code>
	<code>.rel.text</code>
	<code>.symtab</code>
	<code>.bss</code>
	<code>.data</code>
	<code>.rodata</code>
	<code>.text</code>
Program header table (for executables)	
	ELF header

Don't be shy, let's look in an ELF

```
/* Sample program hello.c */
#include <stdlib.h>
#include <stdio.h>

int      OS_EVERYTHING_OKAY = 0;

void      helloWorld      ()
{
    puts("Hello world!");
}

int      main ()
{
    helloWorld();
    return(OS_EVERYTHING_OKAY);
}
```

```
[jphillips@localhost lecture2]$ gcc -c hello.c # makes hello.o
[jphillips@localhost lecture2]$ gcc hello.c -o hello
```

Peering inside an ELF

- ◆ `readelf -h <file>`
 - ◆ Prints header info of `.o .so` and `execs`

```
[jphillips@localhost lecture2]$ readelf -h hello
```

ELF Header:

```

Magic:      7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
Class:                               ELF32
Data:                               2's complement, little endian
Version:                               1 (current)
OS/ABI:                               UNIX - System V
ABI Version:                               0
Type:                               EXEC (Executable file)
Machine:                               Intel 80386
Version:                               0x1
Entry point address:                 0x80483b0
Start of program headers:            52 (bytes into file)
Start of section headers:            2480 (bytes into file)
Flags:                               0x0
Size of this header:                  52 (bytes)
Size of program headers:              32 (bytes)
Number of program headers:            8
Size of section headers:              40 (bytes)
Number of section headers:            30
Section header string table index:    27
```

Peering inside an ELF

◆ `readelf -S <file>`

◆ Prints info on sections

```
[jphillips@localhost lecture2]$ readelf -S hello
```

There are 30 section headers, starting at offset 0x9b0:

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
. . . .										
[20]	.jcr	PROGBITS	08049694	000694	000004	00	WA	0	0	4
[21]	.dynamic	DYNAMIC	08049698	000698	0000e0	08	WA	6	0	4
[22]	.got	PROGBITS	08049778	000778	000004	04	WA	0	0	4
[23]	.got.plt	PROGBITS	0804977c	00077c	00001c	04	WA	0	0	4
[24]	.data	PROGBITS	08049798	000798	000008	00	WA	0	0	4
[25]	.bss	NOBITS	080497a0	0007a0	000008	00	WA	0	0	4
[26]	.comment	PROGBITS	00000000	0007a0	000114	00		0	0	1
[27]	.shstrtab	STRTAB	00000000	0008b4	0000fc	00		0	0	1
[28]	.symtab	SYMTAB	00000000	000e60	000450	10		29	46	4
[29]	.strtab	STRTAB	00000000	0012b0	000250	00		0	0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings)

I (info), L (link order), G (group), x (unknown)

O (extra OS processing required) o (OS specific), p (processor specific)

Peering inside an ELF

- ◆ `readelf -s <file>`
 - ◆ Prints symbol table

```
[jphillips@localhost lecture2]$ readelf -s hello
```

```
Symbol table '.dynsym' contains 7 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	NOTYPE	WEAK	DEFAULT	UND	__gmon_start__
2:	00000000	0	NOTYPE	WEAK	DEFAULT	UND	_Jv_RegisterClasses
3:	00000000	438	FUNC	GLOBAL	DEFAULT	UND	
__libc_start_main@GLIBC_2.0 (2)							
4:	00000000	399	FUNC	GLOBAL	DEFAULT	UND	puts@GLIBC_2.0 (2)
5:	08048394	1020	FUNC	GLOBAL	DEFAULT	UND	
__gxx_personality_v0@CXXABI_1.3 (3)							
6:	08048588	4	OBJECT	GLOBAL	DEFAULT	15	_IO_stdin_used

Symbol table '.symtab' contains 69 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	08048134	0	SECTION	LOCAL	DEFAULT	1	
2:	08048148	0	SECTION	LOCAL	DEFAULT	2	
.

Peering inside an ELF

- `objdump -d -j <section> <file>`
 - Disassembles (“-d”) <section> (“-j”) of <file>

```
[jphillips@localhost lecture2]$ objdump -d -j .text hello.o  
hello.o:          file format elf32-i386
```

Disassembly of section .text:

00000000 <helloWorld>:

0:	55	push	%ebp
1:	89 e5	mov	%esp,%ebp
3:	83 ec 08	sub	\$0x8,%esp
6:	c7 04 24 00 00 00 00	movl	\$0x0,(%esp)
d:	e8 fc ff ff ff	call	e <helloWorld+0xe>
12:	c9	leave	
13:	c3	ret	

00000014 <main>:

14:	8d 4c 24 04	lea	0x4(%esp),%ecx
18:	83 e4 f0	and	\$0xfffffffff0,%esp
1b:	ff 71 fc	pushl	-0x4(%ecx)
1e:	55	push	%ebp
1f:	89 e5	mov	%esp,%ebp
21:	51	push	%ecx
22:	83 ec 04	sub	\$0x4,%esp

. . .

Now we know all about ELFs, *Put the linker to work!*

```
/* one.c */
```

```
#include <stdlib.h>
#include <stdio.h>
```

```
int      start      = 0;
int      stop       = 100;
```

```
extern  int sum ( );
```

```
int      main      ( )
{
    printf("Sum from %d to %d is %d\n",
           start, stop, sum( )
           );
    return(EXIT_SUCCESS);
}
```

◆ Which one.o section has:

◆ main()?

◆ stop?

◆ start?

◆ The string "Sum from %d to %d is %d\n"?

◆ *How can we show this?*

Now we know all about ELF's, *Put the linker to work!*

```
/* two.c */
```

```
extern int start;  
extern int stop;
```

```
int *addrOfStop = &stop;
```

```
int sum()  
{  
    int i;  
    int total = 0;
```

```
    for (i = start; i <= *addrOfStop; i++)  
        total += i;
```

```
    return(total);  
}
```

◆ Which two.o
section has:

◆ `sum()`?

◆ `addrOfStop`?

◆ `i` and `total`?

◆ *Think carefully!*

◆ *How can we show this?*

Linking to make an executable

```
[jphillips@localhost]$ g++ -c one.c  
[jphillips@localhost]$ g++ -c two.c  
[jphillips@localhost]$ g++ -o oneTwo one.o two.o
```

- ◆ Where in oneTwo are:
 - ◆ start and stop? What are their values?
 - ◆ addrOfStop? What is its value?
 - ◆ main() and sum()?

Your turn!

- ◆ Write a program separated into three files:
 - ◆ `first.c`: Has two characters (`begin` and `end`) and `main()` which calls `enterBeginEnd()` and `printFromBeginToEnd()`.
 - ◆ `second.c`: Has `enterBeginEnd()`, which asks for a first character (`begin`), and then asks for a second character (`end`). It must ensure `end` has greater or equal ASCII value.
 - ◆ HINT: `char array[SIZE]; printf("Enter..."); fgets(array, SIZE, stdin); begin = array[0];`
 - ◆ `third.c`: Has `printFromBeginToEnd()` which prints out the characters from `begin` to `end` and their ASCII values in decimal. (Use a simple `for`-loop)

Resolving references

- ♦ The linker ***works*** for its money
 - ♦ It puts the all the `.text`'s together, all the `.data`'s together, *etc.*
 - ♦ It fills in the addresses of pointers, function calls, *etc.*
- ♦ I wonder how we should refer to `jmp` and `call` addresses?
 - ♦ ***Absolute addresses?***
 - ♦ ***Relative addresses?***

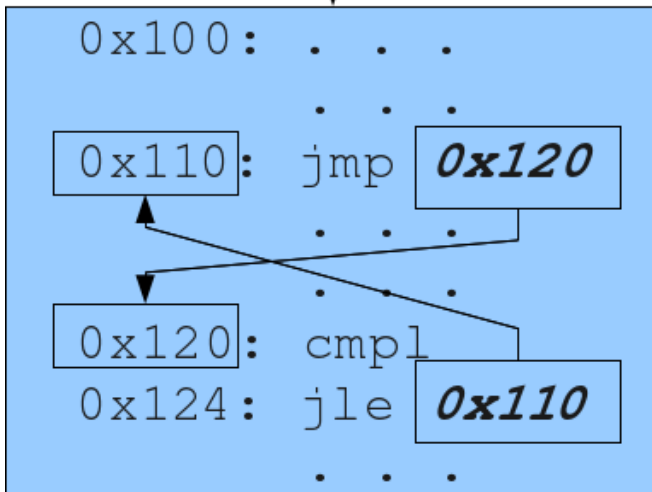
Absolute addressing: compiling Hardcode addresses in `jmp` & `call`

1.c

```
void fnc1 ()
{
    // Look, a loop!
    for ( . . . )
        printf( . . );
}
```

`gcc -c 1.c`

1.o's .text

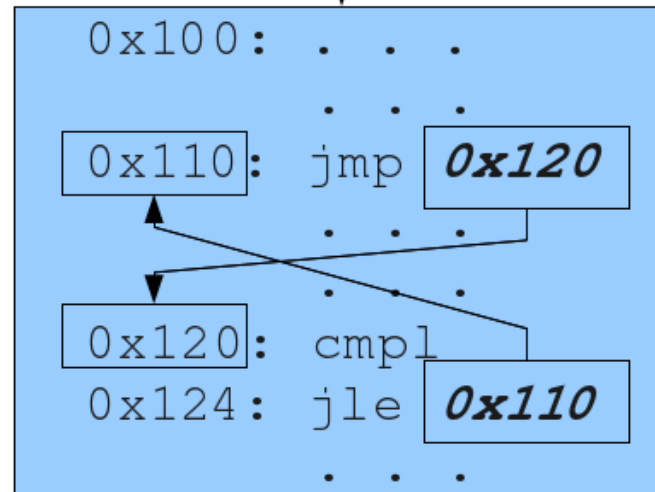


2.c

```
void fnc2 ()
{
    // Another loop!
    for ( . . . )
        printf( . . );
}
```

`gcc -c 2.c`

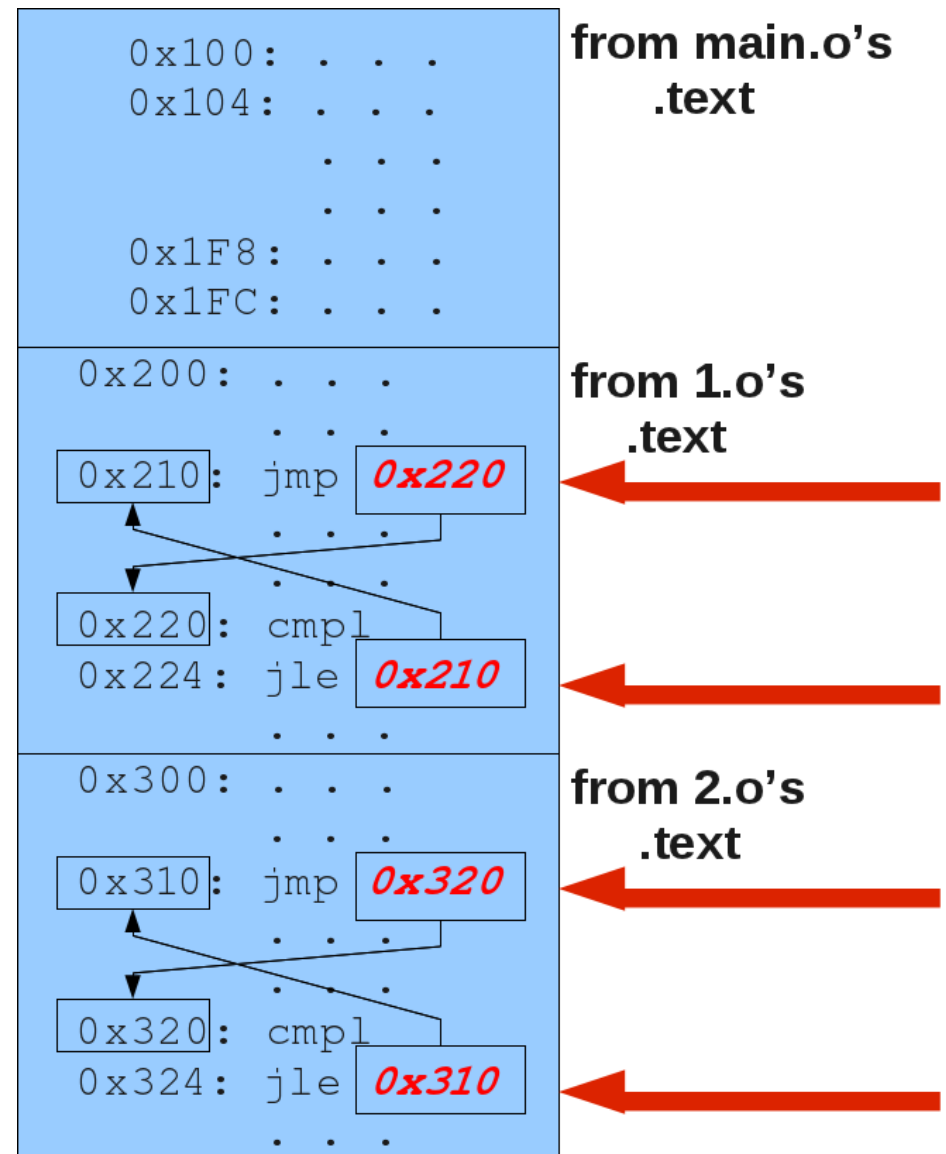
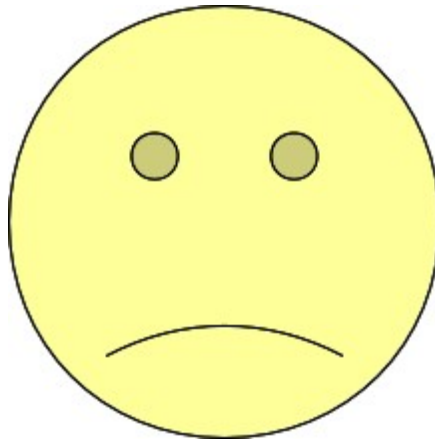
2.o's .text



Absolute addressing: linking

- Uh oh!
 - The linker now has **lots of work** to clean change the addresses of all jmps and calls, **even within same source .texts!**

- Sad linker!



Relative addressing: compiling

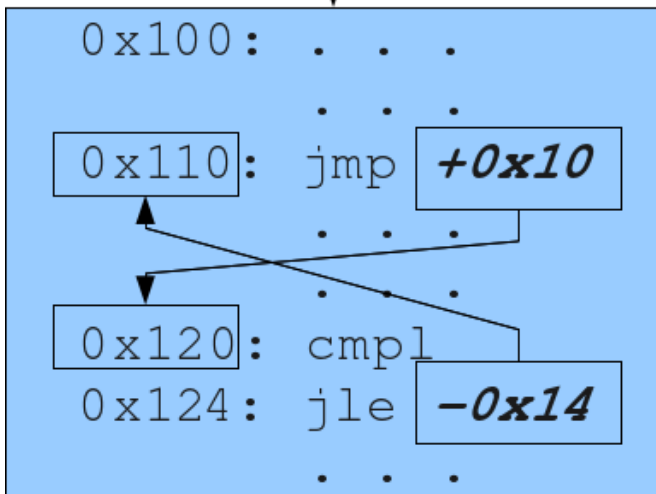
Tell offset (what to add to `eip`) in `jmp` & `call`

1.c

```
void fnc1 ()
{
    // Look, a loop!
    for ( . . . )
        printf( . . );
}
```

gcc -c 1.c

1.o's .text

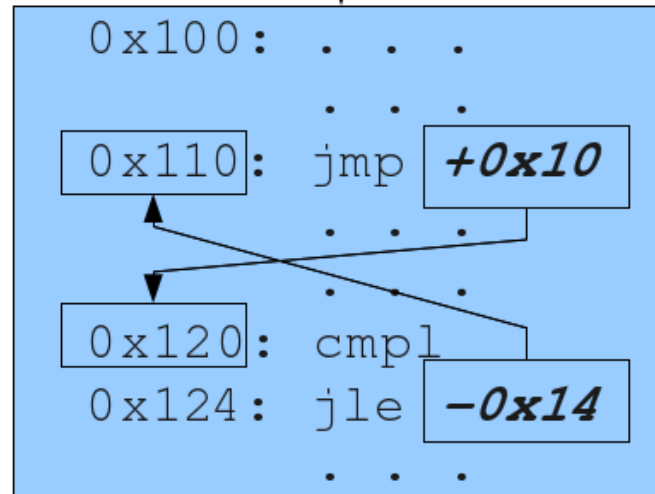


2.c

```
void fnc2 ()
{
    // Another loop!
    for ( . . . )
        printf( . . );
}
```

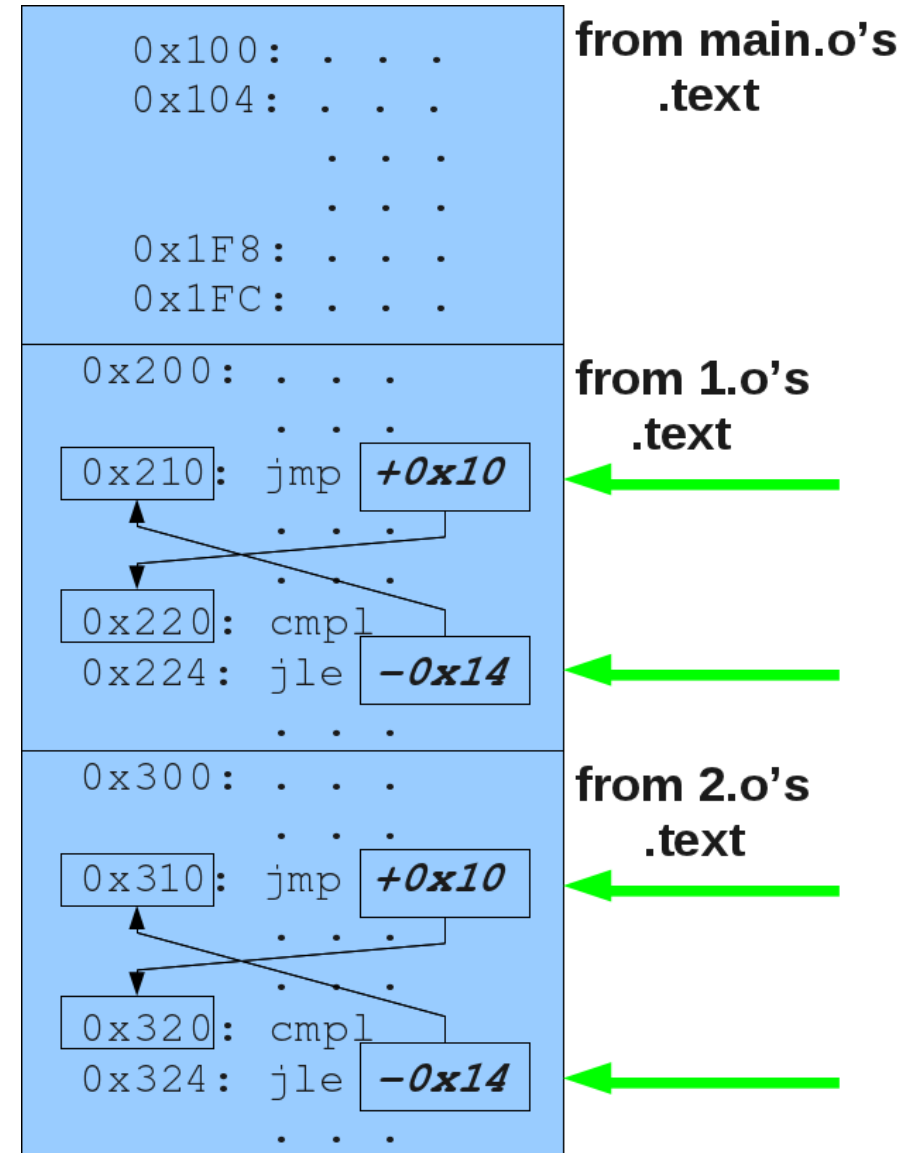
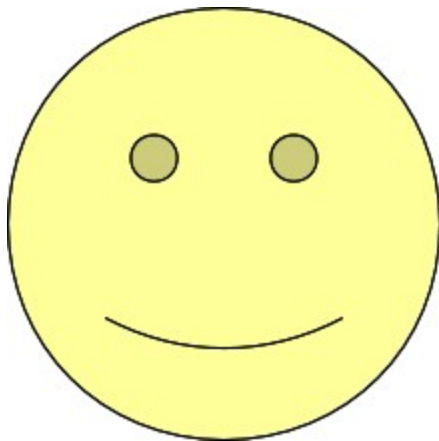
gcc -c 2.c

2.o's .text



Relative addressing: linking

- Easy street!
 - The linker has ***nothing to do*** within same source .texts!
- Happy linker!



Consider jumps in this loop:

```
#include <stdlib.h>
#include <stdio.h>

int main ()
{
    int i;

    for (i = 0; i < 10; i++)
        printf( "%d ", i );

    printf( "\n" );
    return( EXIT_SUCCESS );
}
```

Relative jumps, example 1:

```
080483f4 <main>:
80483f4: 55                push    %ebp
80483f5: 89 e5            mov     %esp,%ebp
80483f7: 83 e4 f0        and     $0xffffffff0,%esp
80483fa: 83 ec 20        sub     $0x20,%esp
80483fd: c7 44 24 1c 00 00 00 movl    $0x0,0x1c(%esp)
8048404: 00
8048405: eb 1a          jmp     8048421 <main+0x2d>
8048407: b8 14 85 04 08   mov     $0x8048514,%eax
804840c: 8b 54 24 1c      mov     0x1c(%esp),%edx
8048410: 89 54 24 04      mov     %edx,0x4(%esp)
8048414: 89 04 24         mov     %eax,(%esp)
8048417: e8 08 ff ff ff   call    8048324 <printf@plt>
804841c: 83 44 24 1c 01   addl    $0x1,0x1c(%esp)
8048421: 83 7c 24 1c 09   cmpl    $0x9,0x1c(%esp)
8048426: 7e df           jle     8048407 <main+0x13>
8048428: c7 04 24 0a 00 00 00 movl    $0xa,(%esp)
804842f: e8 d0 fe ff ff   call    8048304 <putchar@plt>
8048434: b8 00 00 00 00   mov     $0x0,%eax
8048439: c9             leave   %eax
804843a: c3             ret
```

Addr <i>after</i> jmp/call	0x8048407
+ number in jmp/call	+0x000001A
Addr to which to go	0x8048421

You explain this one:

```
080483f4 <main>:
80483f4: 55                push    %ebp
80483f5: 89 e5            mov     %esp,%ebp
80483f7: 83 e4 f0         and     $0xfffffffff0,%esp
80483fa: 83 ec 20         sub     $0x20,%esp
80483fd: c7 44 24 1c 00 00 00 movl    $0x0,0x1c(%esp)
8048404: 00
8048405: eb 1a           jmp     8048421 <main+0x2d>
8048407: b8 14 85 04 08   mov     $0x8048514,%eax
804840c: 8b 54 24 1c     mov     0x1c(%esp),%edx
8048410: 89 54 24 04     mov     %edx,0x4(%esp)
8048414: 89 04 24        mov     %eax,(%esp)
8048417: e8 08 ff ff ff  call    8048324 <printf@plt>
804841c: 83 44 24 1c 01  addl    $0x1,0x1c(%esp)
8048421: 83 7c 24 1c 09  cmpl    $0x9,0x1c(%esp)
8048426: 7e df          jle     8048407 <main+0x13>
8048428: c7 04 24 0a 00 00 00 movl    $0xa,(%esp)
804842f: e8 d0 fe ff ff  call    8048304 <putchar@plt>
8048434: b8 00 00 00 00  mov     $0x0,%eax
8048439: c9             leave
804843a: c3             ret
```

$$\begin{array}{r} 0x8048428 \\ + \quad \quad 0xDF \\ \hline 0x8048407 \end{array}$$

??? How ???

Try explain the
math behind a call
in program you just
wrote

Strong and weak symbols

- ♦ Strong symbols

- ♦ Initialized global vars
- ♦ Functions

```
// Strong  
int initialized = 0;
```

- ♦ Weak symbols:

- ♦ Uninitialized global vars

```
// Weak  
int unInit;
```

- ♦ Rules:

1. Only one strong symbol
2. Weak symbols overridden by strong
3. If all weak linker can choose arbitrarily

```
// Strong  
int someFnc()  
{  
    return(strong+weak);  
}
```

Happy Linker :) or Sad Linker :(?

```
/* 1.c */
```

```
int var;
```

```
int someFnc( )  
{  
    . . .  
}
```

```
/* 2.c */
```

```
int someFnc( )  
{  
    . . .  
}
```

Happy Linker :) or Sad Linker :(?

```
/* 1.c */
```

```
int var;
```

```
int someFnc1()  
{  
    . . .  
}
```

```
/* 2.c */
```

```
int var;
```

```
int someFnc2()  
{  
    . . .  
}
```

Happy Linker :) or Sad Linker :(?

```
/* 1.c */
```

```
int var1;  
int var2;
```

```
int someFnc1()  
{  
    . . .  
}
```

```
/* 2.c */
```

```
double var1;
```

```
int someFnc2()  
{  
    . . .  
}
```


Happy Linker :) or Sad Linker :(?

```
/* 1.c */
```

```
int var1 = 10;
```

```
int someFnc1()  
{  
    . . .  
}
```

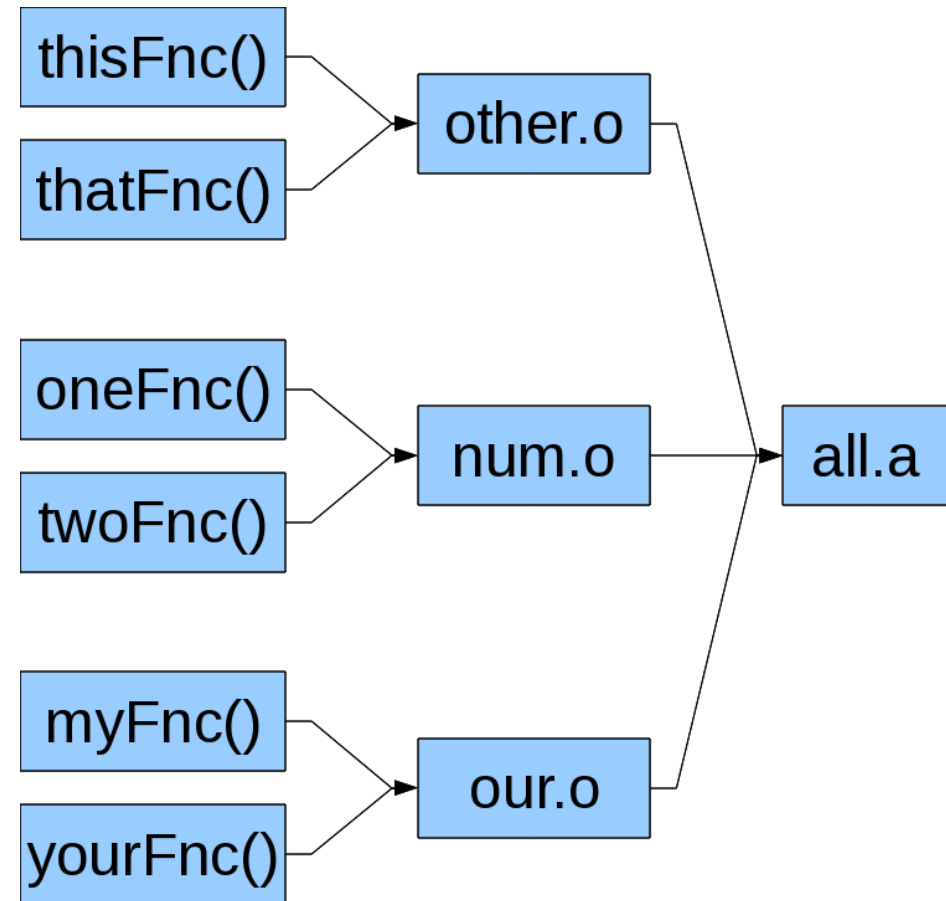
```
/* 2.c */
```

```
double var1;
```

```
int someFnc2()  
{  
    . . .  
}
```

Static libraries (.a for “archives”)

- Link related object files
 - General C functionality
`/usr/lib/libc.a`
 - Math-related (`sin()`, `cos()`, *etc.*)
`/usr/lib/libm.a`
 - Cryptography
`/usr/lib/libcrypt.a`
 - Also for strings, graphics, *etc.*



Static libraries (.a for “archives”)

- ◆ What's in them? Find out:

```
$ ar t /usr/lib/libcrypt.a
```

```
crypt-entry.o
```

```
md5-crypt.o
```

```
md5.o
```

```
sha256-crypt.o
```

```
sha256.o
```

```
sha512-crypt.o
```

```
sha512.o
```

```
crypt.o
```

```
crypt_util.o
```

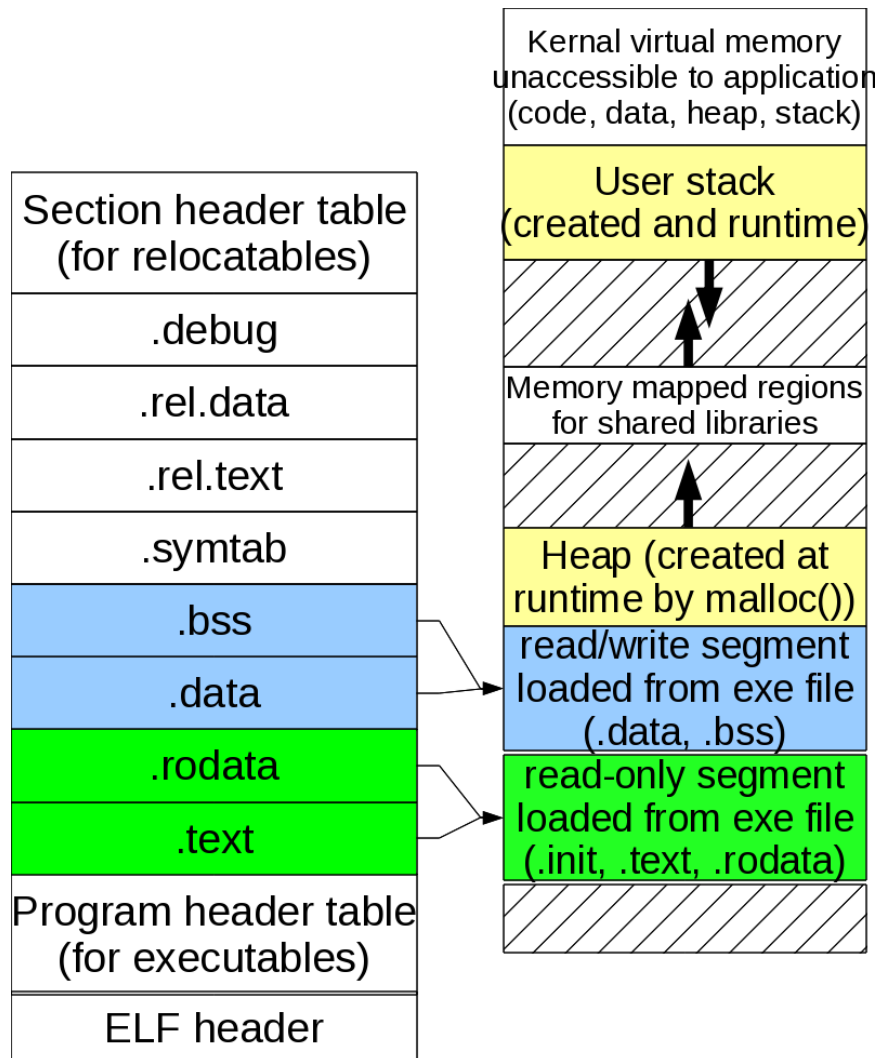
- ◆ How to make them?

```
$ ar rs all.a other.o num.o our.o
```

- ◆ How to use them?

```
$ gcc -o myExecutable main.o -lm
```

Got an executable? *Load it!*



`.text` to execute-only code pages

`.rodata` to read-only data pages

`.data` to read-write data pages

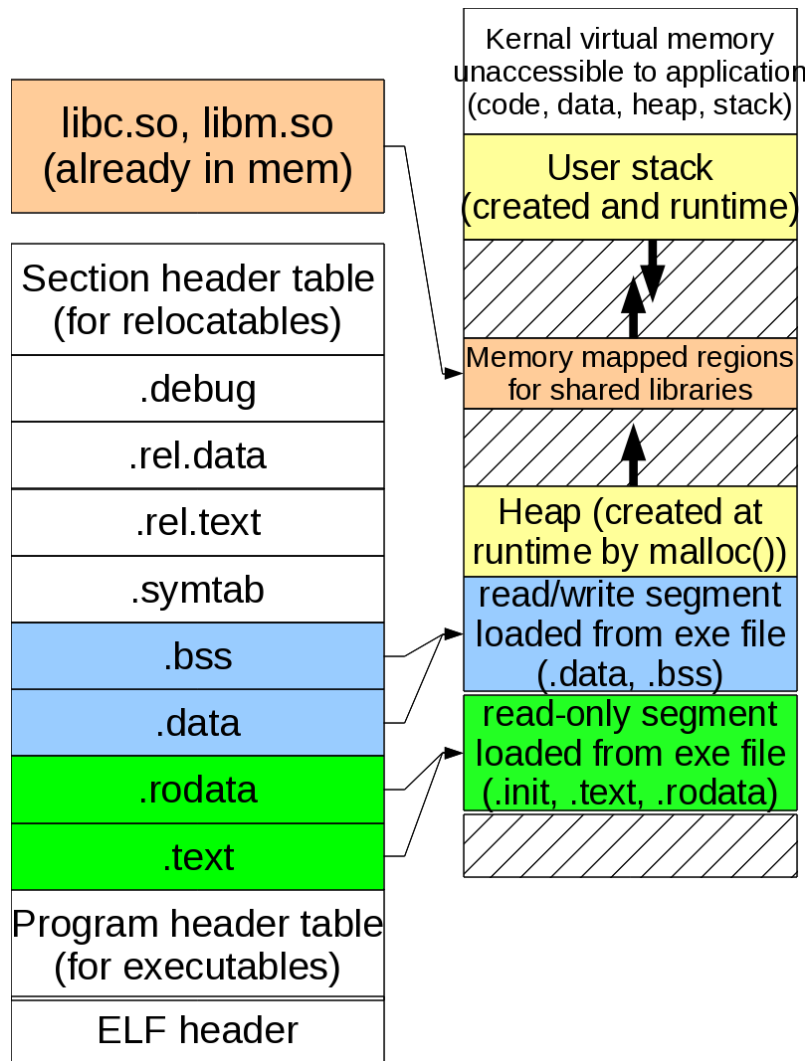
`.bss` expanded to 0-initialized data pages

Construct stack and heap pages

Hey! Aren't we wasting memory?

- ♦ Almost any program written in C includes `libc.a` code for `printf()`, `fgets()`, *etc.*
 - ♦ Wastes **disk-space** (many executables with same code)
 - ♦ (Even worse) wastes **memory** (memory, of course, is faster and more scarce)

Solution: Dynamically Shared Libs



1. `gcc -c myFile.c`
2. `gcc -o myFile.o -lspecial`
3. When loading to execute, dynamically link with `libc.so`, `libm.so`, *etc.*

Don't dynamic linking?

- Use `-static` flag to link `libc.a` instead of dynamically link with `libc.so`.
- But it will cost you though:

```
[jphillips@local]$ gcc -c hello.c
[jphillips@local]$ gcc -o hello hello.o
[jphillips@local]$ gcc -static -o helloStatic hello.c
[jphillips@local]$ ls -l hello helloStatic
-rwxrwxr-x 1      4990 2011-01-11 07:03 hello
-rwxrwxr-x 1  546491 2011-01-11 07:03 helloStatic
```

Next time: *Loading and executing!*

