

CSC 374/407: Computer Systems II

Lecture 7

Joseph Phillips
De Paul University

2014 January 4

Copyright © 2012 Joseph Phillips
All rights reserved

Reading

- ♦ Bryant & O'Hallaron “*Computer Systems, 2nd Ed.*”
 - Chapter 10 (except 10.4): System Level I/O
- ♦ Hoover “*System Programming*”
 - Chapter 5: Input/Output

Topics

High-level C file Input-Output

Iterating over directories

Getting file details

High level C Input-Output

Next lecture will discuss reading and writing ***a buffer of bytes*** efficiently

For now we'll concentrate on the ***high-level approach*** is good for dealing with ***lines, ints, floats, words, etc.***

Uses `FILE*` stream (or `filePtr`) instead of `int` `fileDescriptor`.

Existing `FILE*` files:

- `stdin` (“standard input”)
- `stdout` (“standard output”)
- `stderr` (“standard error”)

fopen()

```
FILE* fopen(const char* pathname,  
            const char* typeP);
```

- Opens file `pathname` according to `typeP`:
- Returns ptr on success or `NULL` otherwise.

`typeP`: can be

- "`r`": reading from beginning
- "`r+`": reading and writing from beginning
- "`w`": writing from beginning (truncated if exists, else create)
- "`w+`": reading and writing from beginning (truncated if exists, created otherwise)
- "`a`": writing from end (create if not exists)
- "`a+`": reading and writing from end (create if not exists)

`fgets()`, `fgetc()`

- ```
char* fgets(char* bufferPtr, int
 bufferLen, FILE* filePtr)
```
- Reads up to `bufferLen-1` characters from `filePtr` into `bufferPtr`. Reads `'\n'` into `buffer` too.
  - Returns `bufferPtr` on success, else `NULL`.
- ```
int fgetc(FILE* filePtr)
```
- Reads up to 1 character from `filePtr`.
 - Returns that char success, else `EOF`.

fprintf()

```
int fprintf(FILE* filePtr, const char*  
format, . . .)
```

- Prints to substituted format to filePtr.
- Substitutions include:
 - %d: Substitute in integer as decimal number
 - %x, %X: Substitute in integer as hexadecimal number
 - %c: Substitute in character
 - %s: Substitute in string
 - %g, %f: Substitute in floats and doubles
 - %p: Substitute in pointer value
- Returns returns number chars printed.
- printf() is the same as fprintf(stdout, . . .)

`fflush()`, `fclose()`

```
int fflush(FILE* filePtr)
```

- Flushes `filePtr` to disk, screen *etc.*
- `fflush(stdout)`:
 - Works fine in Linux,
 - May be problematic in Microsoft C.
- Returns 0 on success, otherwise `errno` is set.

```
fclose(FILE* filePtr)
```

- Closes `filePtr`.
- Returns 0 on success, otherwise `errno` is set.

Your turn!

Write a program that takes two parameters:

`$lineCounter string filename`

that counts and returns the number of lines of
`filename` that begin with `string`.

- If `filename` cannot be opened it writes an error message to `stderr`.

Well, there is `fscanf()`, but . . .

Just so you've seen it:

```
int fscanf(FILE* filePtr, const char*  
    format, . . . )
```

- Returns number of items read

Better to use `fgets()`, then

```
int sscanf(const char* source, const  
    char* format, . . . );
```

- What goes in format? Largely the same codes as for `fprintf()` (next slide).

```
int strtol(const char*, char**, int)
```

- Returns integer: `strtol("123", 0, 10) == 123`

```
double strtod(const char*, char**)
```

- Returns double: `strtod("12.3", 0) == 12.3`

Like FILE* but want buffered objects instead of lines?

```
size_t fread(void* ptr, size_t size,  
size_t numItems, FILE* filePtr)
```

- Reads numItems of size size from filePtr and puts them in ptr.
- Returns number items read.

```
size_t fwrite(const void* ptr, size_t  
size, size_t numItems, FILE* filePtr)
```

- Writes numItems of size size from ptr to filePtr.
- Returns number items written.

Your turn!

Write a program that reads from 0 to N int pairs:

- Ignore blank lines or lines with just spaces
- Ignore lines whose first non-space char is # as comments
- Ignore any spaces up to the two ints, and between them
- Uncommented letters, *etc.* are errors.

Eeww! Parsing!

- What's the best programming structure to read an unbounded number of lines?
- Useful stuff:
 - `int isdigit(char c), int isspace(char c)`

stdout VS. stderr

Q: Why might it be useful to distinguish between output messages and error messages?

A: For debugging!

```
#include <stdlib.h>
#include <stdio.h>
/* $ ./stdoutVsStderr
 * I'm an ordinary msg.
 * I'm the error msg.
 * $ ./stdoutVsStderr 2> error.txt
 * I'm an ordinary msg.
 * $ cat error.txt
 * I'm the error msg.
 */
int main ()
{
    fprintf(stdout, "I'm an ordinary msg.\n");
    fprintf(stderr, "I'm the error msg.\n");
    return(EXIT_SUCCESS);
}
```

Is using `FILE*` as efficient as `int fd`?

Probably not (`FILE*` uses `int fd`), but it is buffered.

```
#include <stdlib.h>
#include <stdio.h>
```

```
int main()
{
    printf("T");
    printf("h");
    printf("i");
    printf("s");
    printf(" ");
    printf("i");
    printf("s");
```

```
    printf("\n");
    printf("'");
    printf("t");
    printf(" ");
    printf("e");
    printf("f");
    printf("f");
    printf("i");
    printf("c");
    printf("i");
    printf("e");
    printf("n");
    printf("t");
    printf("\n");
    fflush(stdout);
    return(EXIT_SUCCESS);
}
```

Is using FILE* as efficient as int fd?

```
$ strace ./printf_sys_call_ex
execve("./printf_sys_call_ex", [ "./printf_sys_call_ex" ],
    [ /* 46 vars */ ]) = 0
brk(0)                                = 0x8fa7000
access("/etc/ld.so.preload", R_OK)    = -1 ENOENT
open("/etc/ld.so.cache", O_RDONLY)    = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=63949, ...})
                                         = 0
mmap2(NULL, 63949, PROT_READ, MAP_PRIVATE, 3, 0)
                                         = 0xb7fb3000
close(3)                               = 0
open("/lib/libc.so.6", O_RDONLY)      = 3
read(3,
    "\177ELF\1\1\1\0\0\0\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\360\364
    @\0004\0\0\0"..., 512)             = 512
. . .
write(1, "This isn't efficient\n", 21This isn't efficient
)                                       = 21
```

Manipulating files and filesystems

There are several other system calls for the Unix file system including:

```
#include <unistd.h>
```

```
#include <sys/stat.h>
```

```
unlink(const char* filename);
```

- Removes (erases) files.

```
chmod(const char *path, mode_t mode);
```

- Changes file permissions

```
chdir(const char *path);
```

- Changes the working directory

Iterating over files in directory

Like `fopen()`, `fgets()`, `fclose()` but for directories

```
#include <sys/types.h>
#include <dirent.h>
```

```
DIR*          opendir   (const char* name);
struct dirent* readdir   (DIR *dir);
int           closedir  (DIR*);
```

```
struct dirent
{
    ino_t    d_ino;        // inode number
    off_t    d_off;        // offset to next dirent
    ushort   d_reclen;     // length of record
    uchar     d_type;       // type of file
    char      d_name[256];  // filename
};
```

Your turn!

Write a program `lister` that takes an optional command line argument

- `./lister dirName`
 - Lists directory `dirName` (assume it exists)
- `./lister badDirName`
 - Prints an error message to `stderr` if `badDirName` is not a directory or if don't have permission to read it.
- `./lister`
 - Lists the items in the current directory (“.”)

Finding details about a file:

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
int stat(const char* path, struct stat* buf);
struct stat
{
    dev_t      st_dev;      // Device ID
    ino_t      st_ino;      // inode
    mode_t     st_mode;     // what type of "file"
    nlink_t    st_nlink;    // num hard links
    uid_t      st_uid;      // user ID of owner
    gid_t      st_gid;      // group id of owner
    dev_t      st_rdev;     // Device ID (special files)
    off_t      st_size;     // Total size in bytes
    blksize_t  st_blksize;  // Filesys' block size
    blkcnt_t   st_blocks;   // Num allocated blocks
    time_t     st_atime, st_mtime, st_ctime;
    // Access (read or write), modify (change metadata),
    // change (write) times
};
```

stat, cont'd

What type of file is that?

Use these macros on `st_mode`:

- `S_ISREG(m)`: Regular file
- `S_ISDIR(m)`: Directory
- There are others (block & char devices, symbolic links, FIFOs and sockets)

Your turn!

Revise your `lister` program into `lister2` that for files will print:

- the size in bytes for files
- "`dir`" for directories
- "`other`" of entries other than a file or directory

stat, cont'd

“Hey buddy, got the time?” Recall:

```
struct stat
{
    . . .
    time_t    st_atime;    // Last Access (read or write)
    time_t    st_mtime;    // Last Modify (metadata)
    time_t    st_ctime;    // Last Change (write)
};
```

Printing the time:

```
#include <time.h>
```

```
char*  ctime(time_t*  );
```

- Returns c-string telling time in human-readable form

Your turn!

Revise `lister2` to print the last change (write) time for all entries

How would you modify your program to recursively descend into directories (other than “.” and “..”)

Next time: Low-level I/O and Sockets