

# CSC 374/407: Computer Systems II

## Lecture 8

Joseph Phillips  
De Paul University

2014 January 5

Copyright © 2012 Joseph Phillips  
All rights reserved

# Reading

- ♦ Bryant & O'Hallaron “*Computer Systems, 2<sup>nd</sup> Ed.*”
  - Chapter 10 (except 10.4): System Level I/O
  - Chapter 11: Networking Programming
- ♦ Hoover “*System Programming*”
  - Chapter 5: Input/Output

# Topics

Unix Filesystem Design: A Process' Prospective

Unix Filesystem Design: A Systemwide Prospective

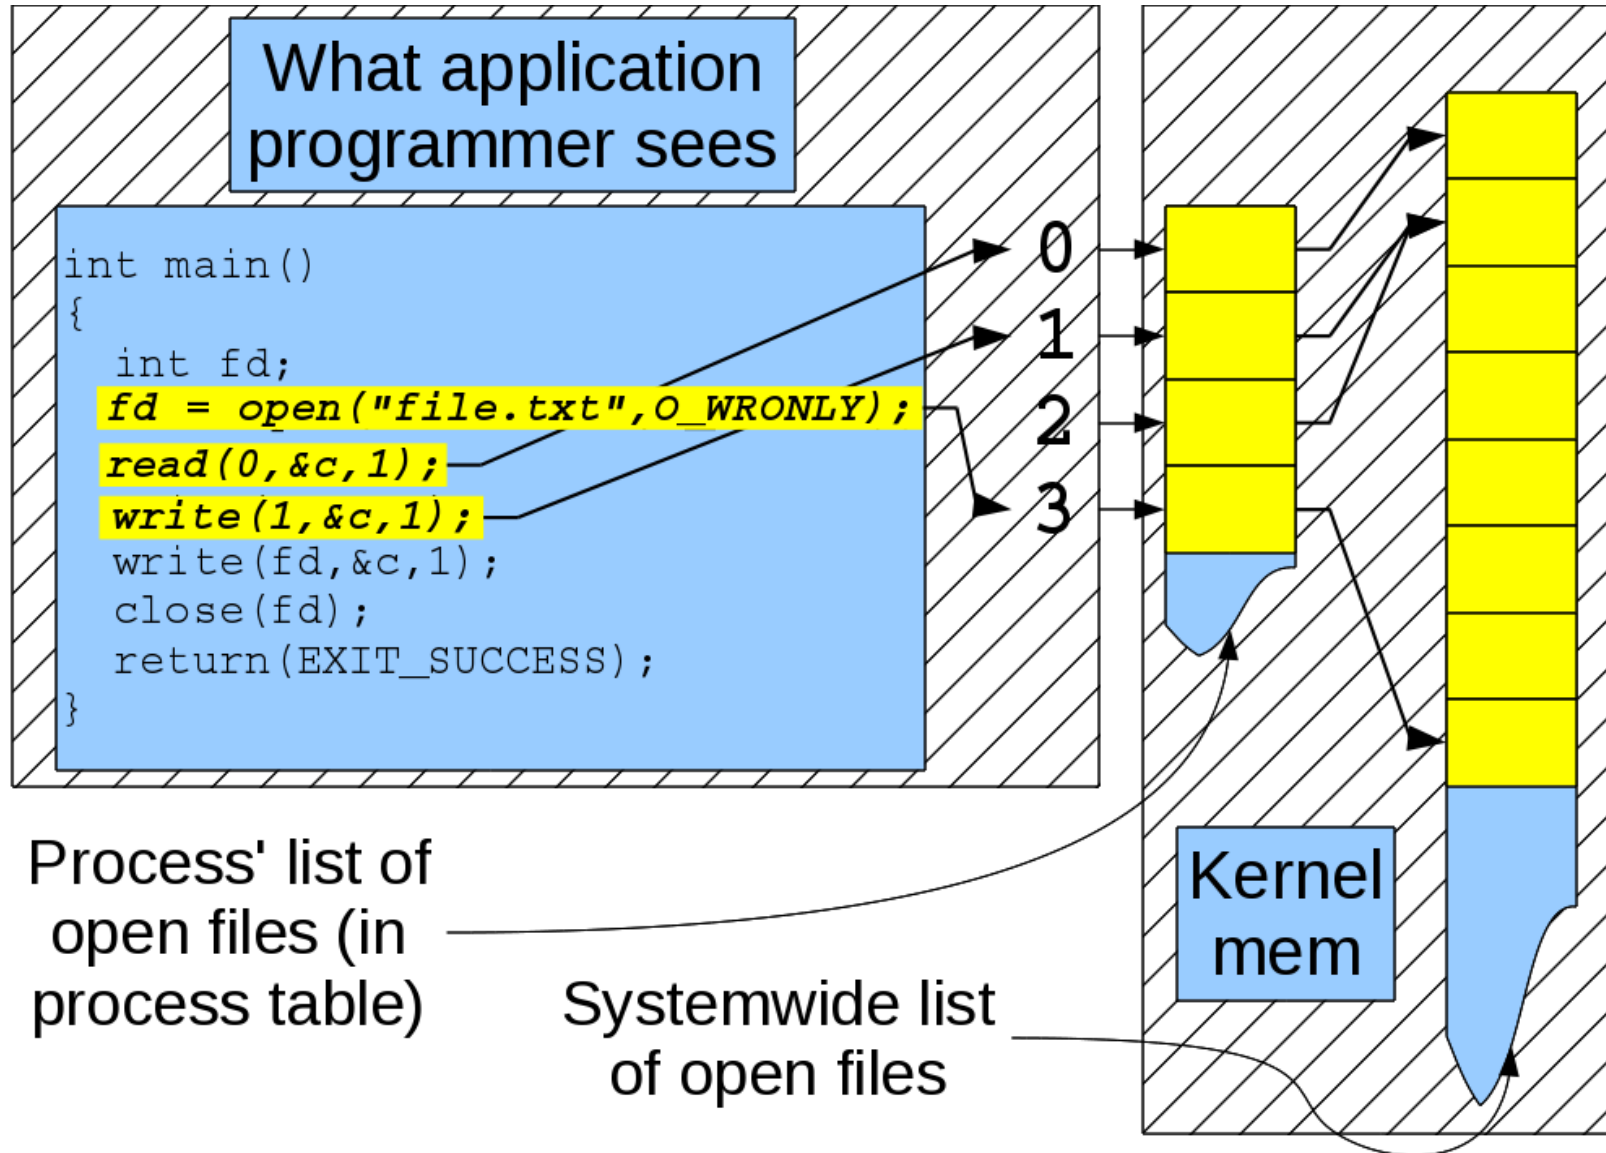
Low-level C Input-Output

Socket communication and the client/server model

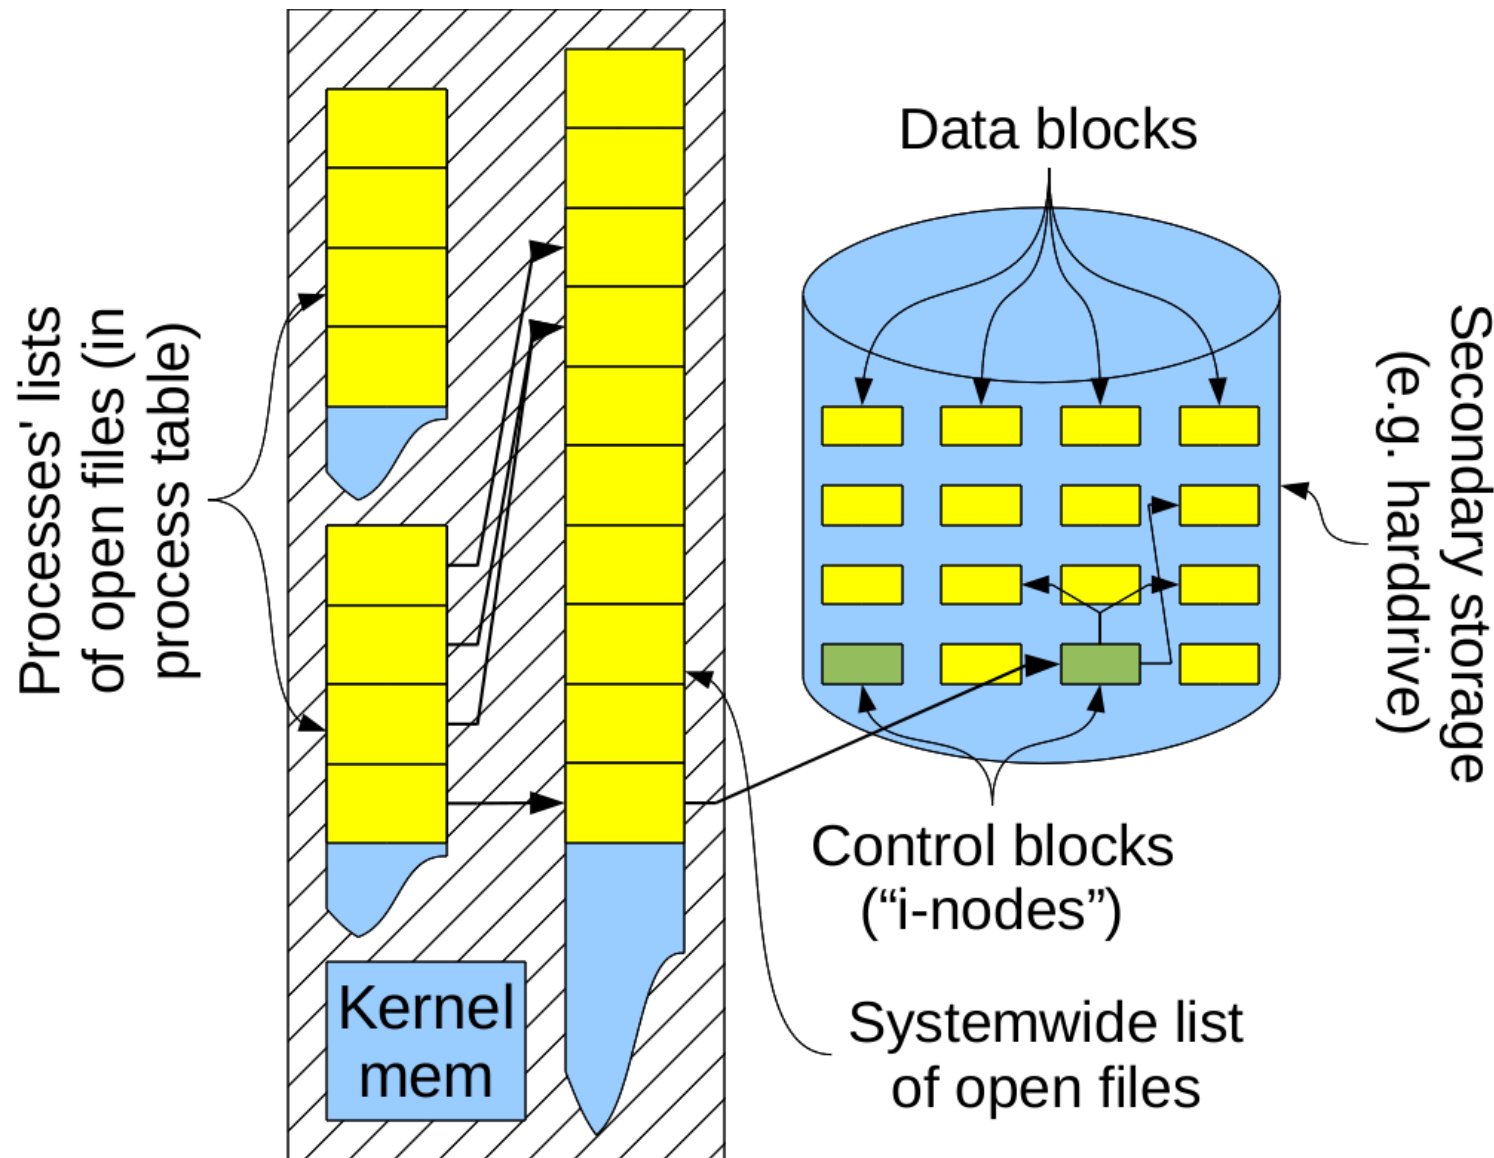
Server-side socket programming

Client-side socket programming

# Unix Filesystem Design: A Process' Prospective



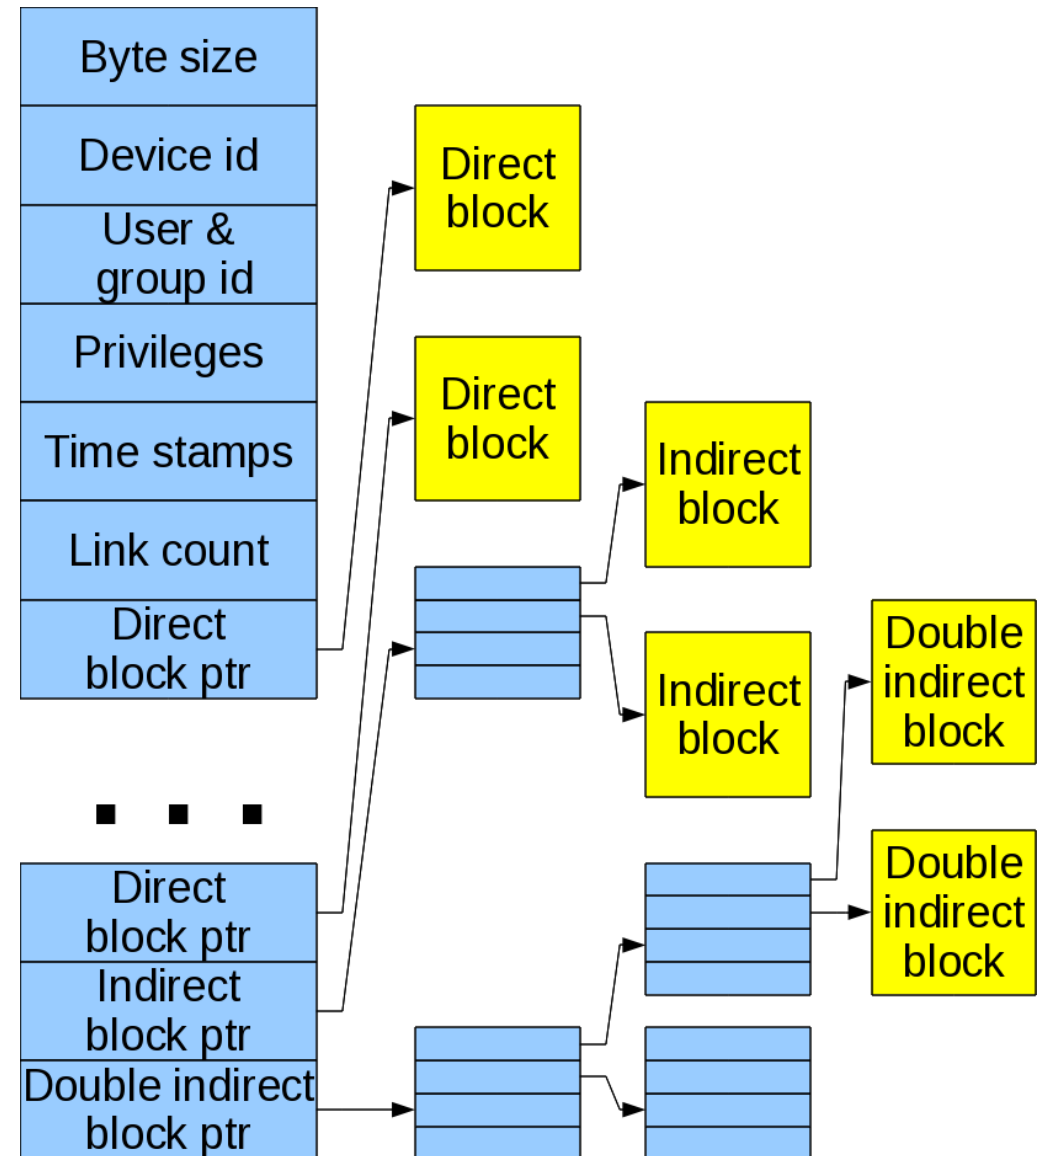
# Unix Filesystem Design: A Systemwide Prospective



# What's an “I-Node”?

Tells a files:

- Size in bytes
- Access times (last read, last written, last its status was modified)
- User and group ID
- Device ID
- Access privileges
- Link count (num different names/directories)
- Pointers



# Low level C Input-Output

File descriptors are indices into process' file table

- 0: Standard input (`stdin`)
- 1: Standard output (`stdout`)
- 2: Standard error (`stderr`)

Useful commands include:

```
int open(const char* path, int how, int
    permission)
```

```
int close(int fd)
```

```
int read(int fd, char* bufferPtr,
    size_t bufferSize)
```

```
int write(int fd, char* bufferPtr,
    size_t numBytes)
```

```
int dup(int fd);
```

```
int pipe(int** );
```

# open ( )

```
int open(const char* path, int how,  
int permission)
```

- Returns file descriptor (index into process' file array)
- File path given by `path`.
- Integer `how` is bitwise or-ing of one of:
  - `O_RDONLY`: Open for reading only.
  - `O_WRONLY`: Open for writing only.
  - `O_RDWR`: Open for reading and writing.
- And one or more of:
  - `O_CREAT`: Create file if doesn't already exist
  - `O_TRUNC`: If exist truncate its length to 0 (even if not open for writing)
  - `O_EXCL`: If `O_CREAT` is also set fail if file exists.
  - `O_APPEND`: Write to end of file.



# `open()`, `cont'd`

```
int open(const char* path, int flags,  
int permission)
```

– Permissions are the bitwise or-ing (|) of:

- $1 \ll 8 = 0x100 = 0400$  Owner read
- $1 \ll 7 = 0x80 = 0200$  Owner write
- $1 \ll 6 = 0x40 = 0100$ : Owner execute
- $1 \ll 5 = 0x20 = 0040$ : Group read
- $1 \ll 4 = 0x10 = 0020$ : Group write
- $1 \ll 3 = 0x8 = 0010$ : Group execute
- $1 \ll 2 = 0x4 = 0004$ : Others read
- $1 \ll 1 = 0x2 = 0002$ : Others write
- $1 \ll 0 = 0x1 = 0001$ : Others execute

# close( )

```
int close (int fd)
```

- Closes fd.
- Returns 0 on success or -1 otherwise.
- Does not flush file.

# Your turn!

You are going to open a process' first file after standard input (0), standard output (1) and standard error (2). The file descriptor is an index in a table. What is its integer value?

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>

int main ()
{
    int fd = open("bubu.txt", O_WRONLY | O_CREAT | O_APPEND, 0660);
    printf("fd = %d\n", fd);
    close(fd);
    return(EXIT_SUCCESS);
}
```

# **write( )**

```
int write(int fd, char* bufferPtr,  
size_t numBytes)
```

- Writes numBytes pointed to by bufferPtr to fd.
- Returns number of bytes written, or -1 on error.

# write ( ) example

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>

int main ( )
{
    int fd = open( "bubu.txt" , O_WRONLY | O_CREAT | O_APPEND , 0660 ) ;
    write( fd , "Bubu!\n" , 6 ) ;
    close( fd ) ;
    system( "ls -l ./bubu.txt" ) ;
    return( EXIT_SUCCESS ) ;
}
```

# read( )

```
int read(int fd, char* bufferPtr,  
size_t bufferSize)
```

- Reads up to `bufferSize` bytes from `fd` and puts them into `bufferPtr`.
- Returns number of bytes read from file, either
  - 0 (*“No more left!”*),
  - `bufferSize` (*“Here's a whole buffer full!”*),
  - somewhere inbetween (*“Here's all that's left”*), or,
  - -1 (*“Error!”*)

# read( ) example

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>

#define BUFFER_SIZE 256

int main ( )
{
    char buffer[BUFFER_SIZE];
    int  fd = open( "bubu.txt", O_RDONLY, 0660 );

    read( fd, buffer, BUFFER_SIZE );
    printf( "%s", buffer );
    close( fd );
    return( EXIT_SUCCESS );
}
```

# Your turn!

Write your own simple version of the Unix `littleCopy` file copying command. I'll get you started:

```
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

#define          BUFFER_SIZE          256

/* Continued on next slide */
```



# Your turn!

```
/* From previous slide */

int main (int argc, const char* argv[])
{
    const char* fromFileCPtr;
    const char* toFileCPtr;

    if (argc < 3)
    {
        fprintf(stderr,
                "Usage: littleCopy <fromFile> <toFile>\n"
                );
        return(EXIT_FAILURE);
    }

    fromFileCPtr = argv[1];
    toFileCPtr   = argv[2];

    /* YOUR CODE HERE */
    return(EXIT_SUCCESS);
}
```

# Your turn, again!

```
/* Write a program that counts the number of occurrences
   of a character given on the command line. */
int main (int argc, const char* argv[])
{
    const char  charToCount;
    const char* fileCPtr;

    if (argc < 3)
    {
        fprintf(stderr,
                "Usage: charCount <char> <file>\n"
                );
        return(EXIT_FAILURE);
    }

    charToCount = *argv[1];
    fileCPtr    = argv[2];

    /* YOUR CODE HERE */
    return(EXIT_SUCCESS);
}
```

# **And your turn, yet again!**

Revise the previous program to  
count the number of lines in a file.

# What happens if mother and child write to same file?

```
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

#define FILENAME "bubu.txt"

int main ()
{
    const char* wordsPtr;
    int i;
    int numBytes;
    int pid;
    int fd =
open(FILENAME,
    O_WRONLY | O_CREAT | O_TRUNC,
    0660);

    if (fd < 0)
    {
        fprintf(stderr,
            "Sorry, I can't make "
            "the output file %s\n",
            FILENAME);
        return(EXIT_FAILURE);
    }

    pid = fork();

    if (pid < 0)
    {
        fprintf(stderr,
            "Too many processes ace!\n"
            );
        return(EXIT_FAILURE);
    }
```

# What happens if mother and child write to same file?

```
else
if  (pid == 0)
    wordsPtr =
        "Baby says \"Gaga Gugu!\"\\n";
else
    wordsPtr =
        "Mama says \"Poor baby!\"\\n";

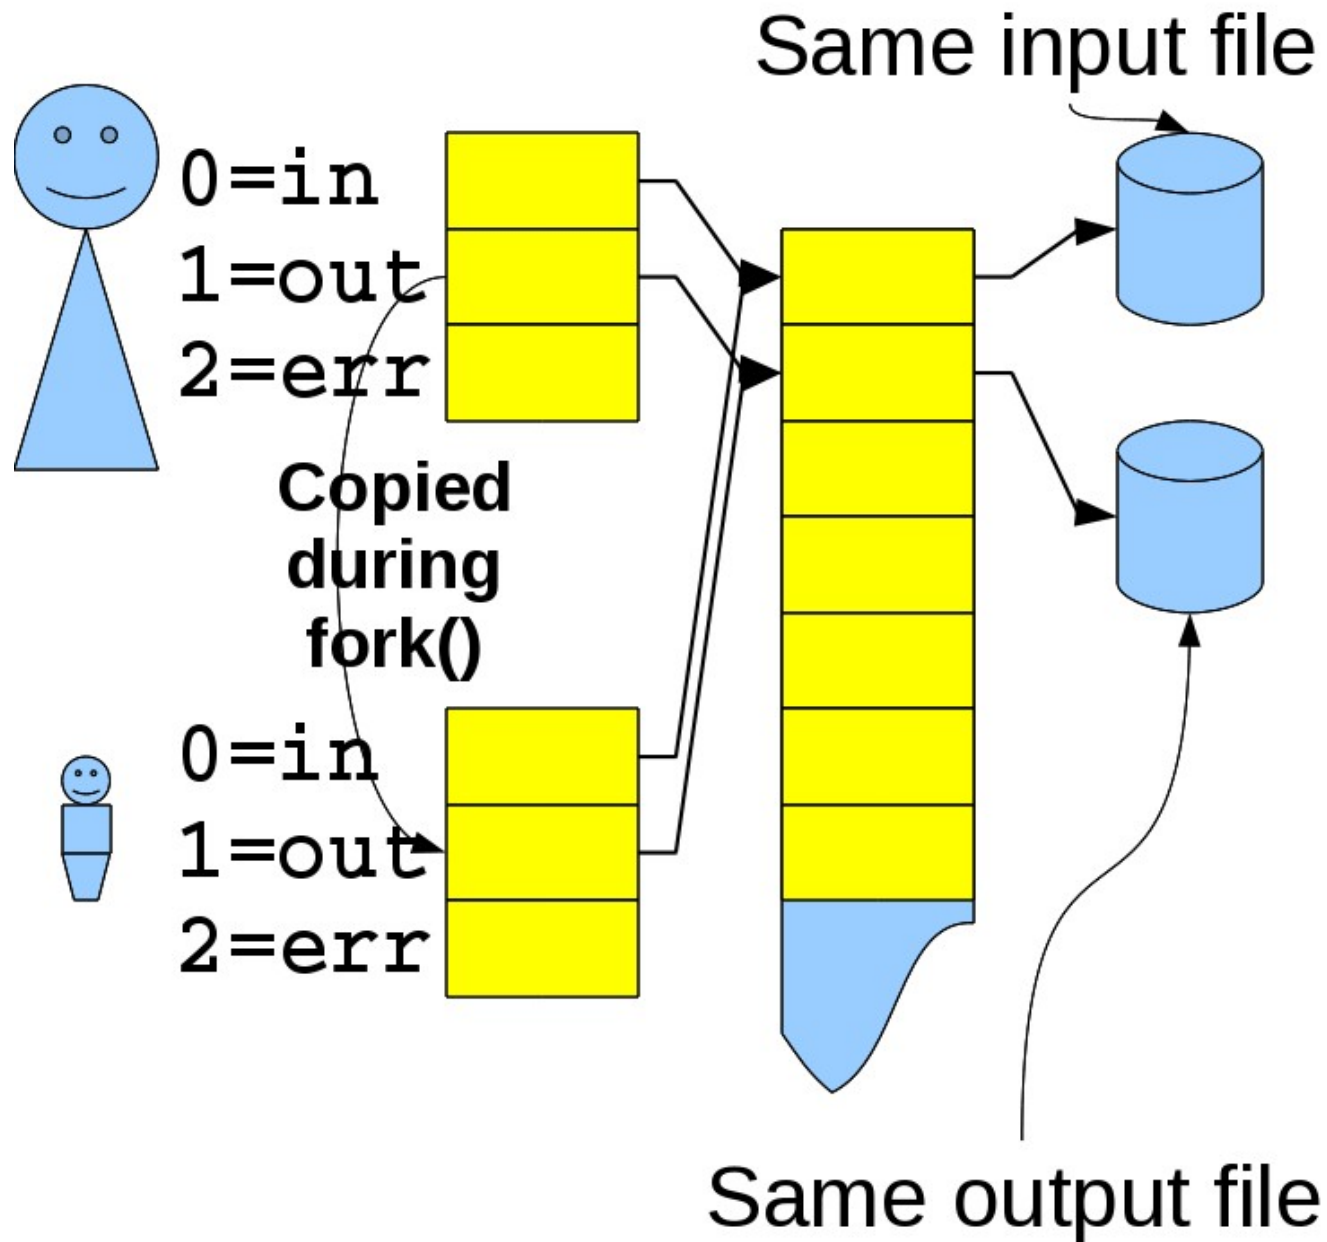
numBytes = strlen(wordsPtr);

for  (i = 0;  i < 4;  i++)
{
    sleep(1);
    write(fd, wordsPtr,
          numBytes);
    printf(wordsPtr);
}

if  (pid > 0)
{
    sleep(2);
    close(fd);
}

return(EXIT_SUCCESS);
}
```

# What's going on?



# Hey! Maybe we can use this for interprocess communication!

```
#include <unistd.h>
. . .
const int PIPE_READ = 0;
const int PIPE_WRITE = 1;
int      myPipe[2];

if (pipe(myPipe) == 0)
{
    char myArray[6];
    write(myPipe[PIPE_WRITE], "Hello!", 6);
    read (myPipe[PIPE_READ ], myArray, 6);
}
```



myPipe: An OS-owned buffer

# dup()

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
#define FILENAME    "bubu.txt"
```

```
int main ()
{
    in fd= open(FILENAME,
                O_WRONLY|O_CREAT|O_TRUNC,
                0660);

    close(1); // Close stdout
    dup(fd);  // Redirect stdout to FILENAME
    printf("I wonder where this will show up?\n");
    close(fd); // Be polite!
    return(EXIT_SUCCESS);
}
```

stdin	0
stdout	1
stderr	2
<b>bubu.txt</b>	3

stdin	0
<del>stdout</del>	1
stderr	2
bubu.txt	3

stdin	0
<b>bubu.txt</b>	1
stderr	2
bubu.txt	3

**dup ( )** copies the entry of the given file descriptor to the first free one.

1

2

3

1



# IPC with pipes

```
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

int      main      ( )
{
    int    parentToChild[2];
    int    childToParent[2];

    if    ( (pipe(parentToChild) < 0)
            || (pipe(childToParent) < 0) )
    {
        fprintf(stderr,
                "Can't make pipes\n");
        return(EXIT_FAILURE);
    }

    if    (pid < 0)
    {
        fprintf(stderr, "Too many
processes Ace!\n");
        return(EXIT_FAILURE);
    }
    else
    if    (pid == 0)
    {
        //  Baby's case
        close(0);    // Close "stdin"
        dup(parentToChild[0]);
        close(1);    // Close "stdout"
        dup(childToParent[1]);

        // . . . continued

        int    pid      = fork();
```

# IPC with pipes, cont'd

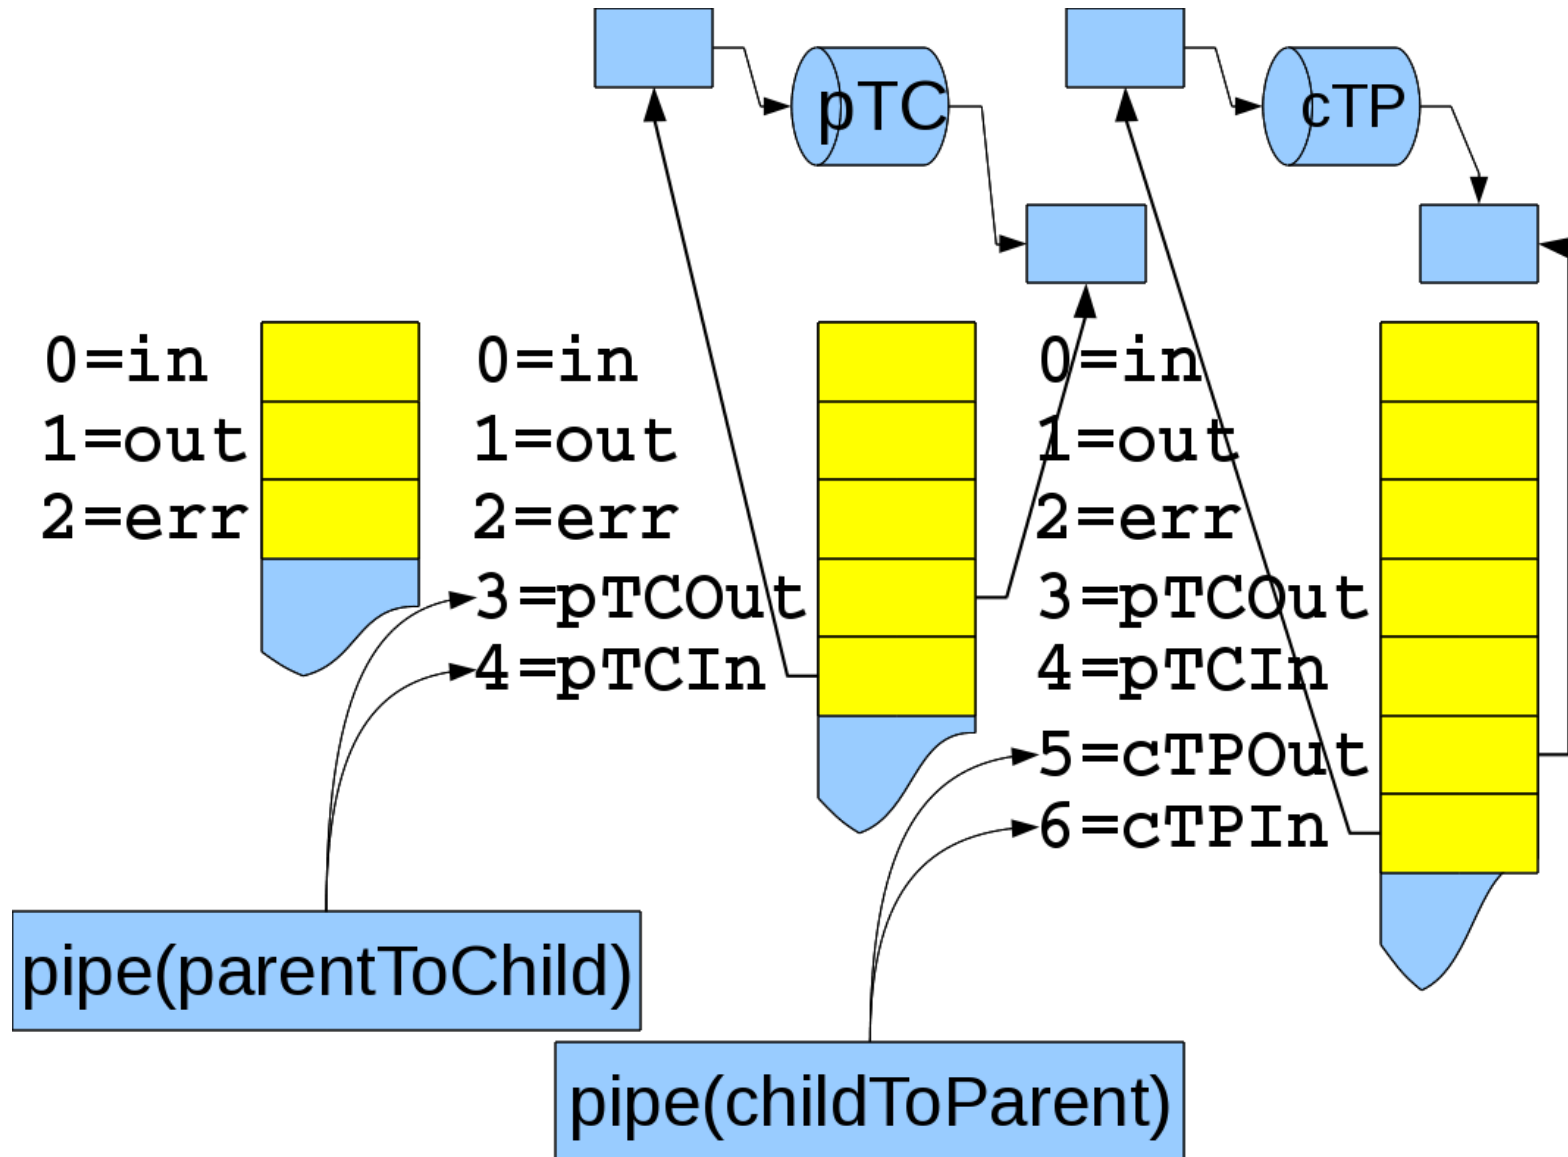
```
// Baby's case, continued
while (1)
{
    char buffer[10];
    int i,numRead;
    numRead =
        read(0,buffer,10);

    for (i=0; i<numRead; i++)
        buffer[i] =
            toupper(buffer[i]);

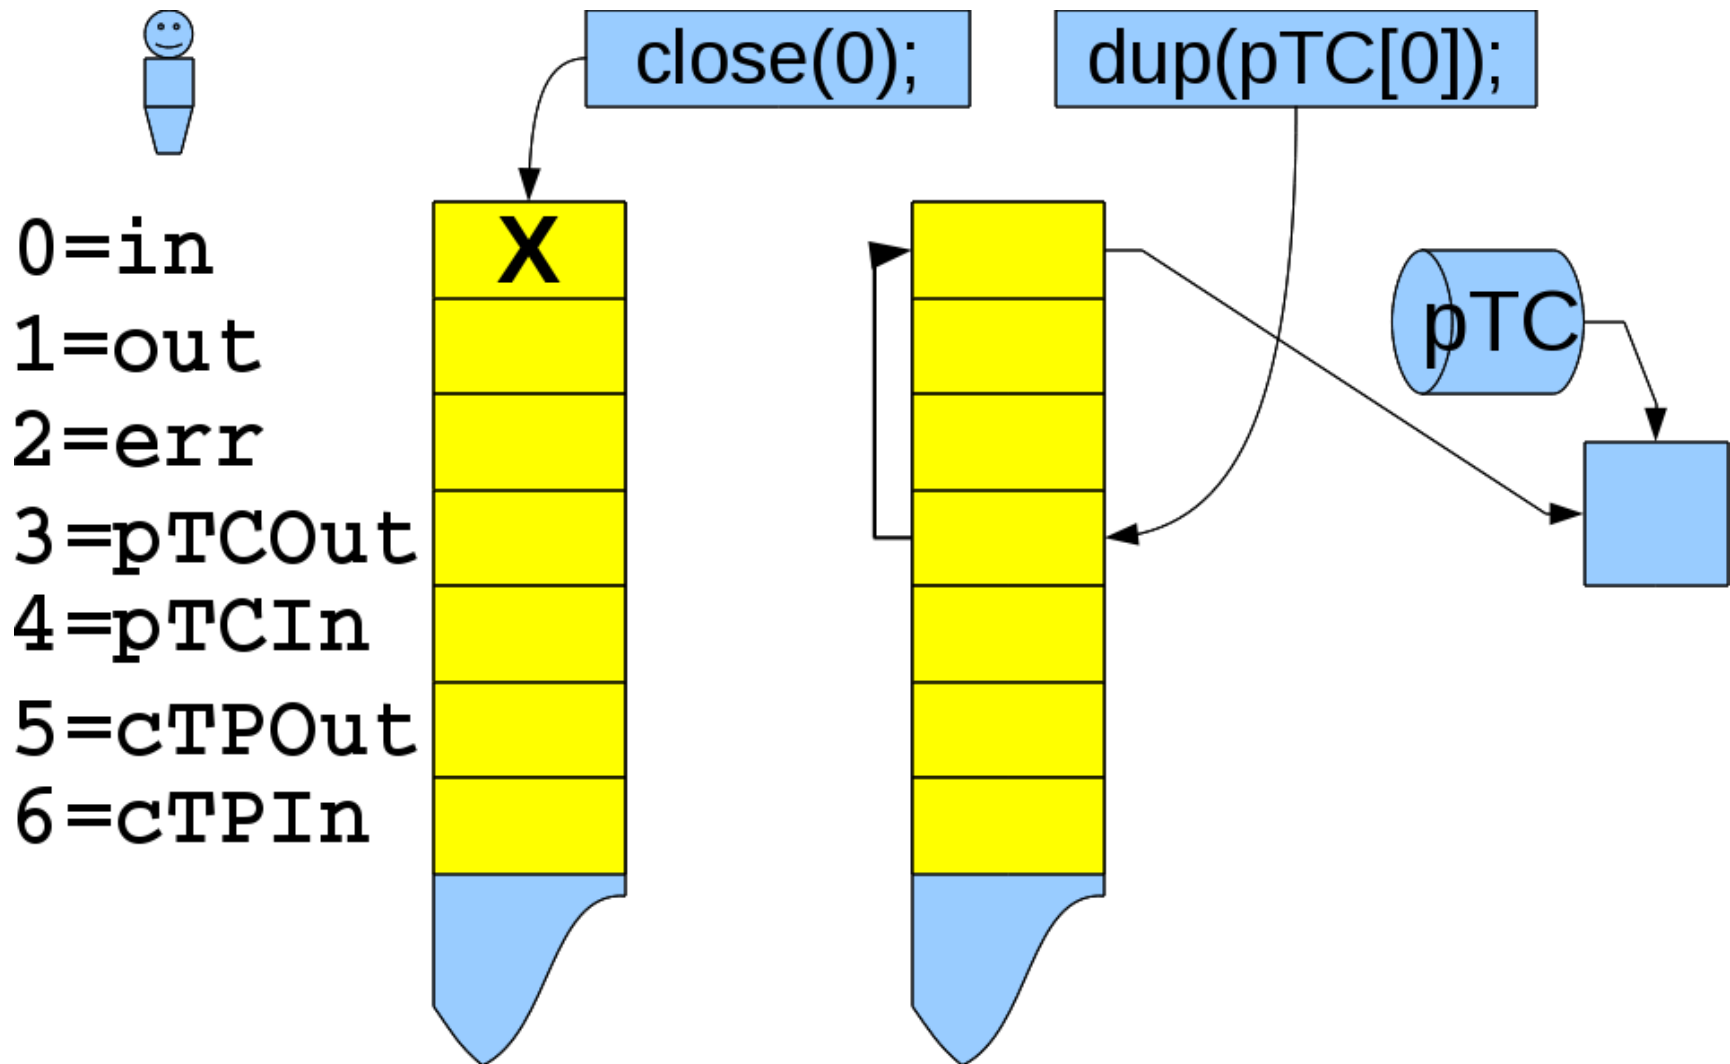
    write(1,buffer,numRead);
}
}
else
{
    // Mama's case
    while (1)
    {
        char buffer[10];
        fgets(buffer,10,stdin);
        write
            (parentToChild[1],
             buffer,
             10);

        read
            (childToParent[0],
             buffer,
             10);
        printf(buffer);
    }
}
return(EXIT_SUCCESS);
}
```

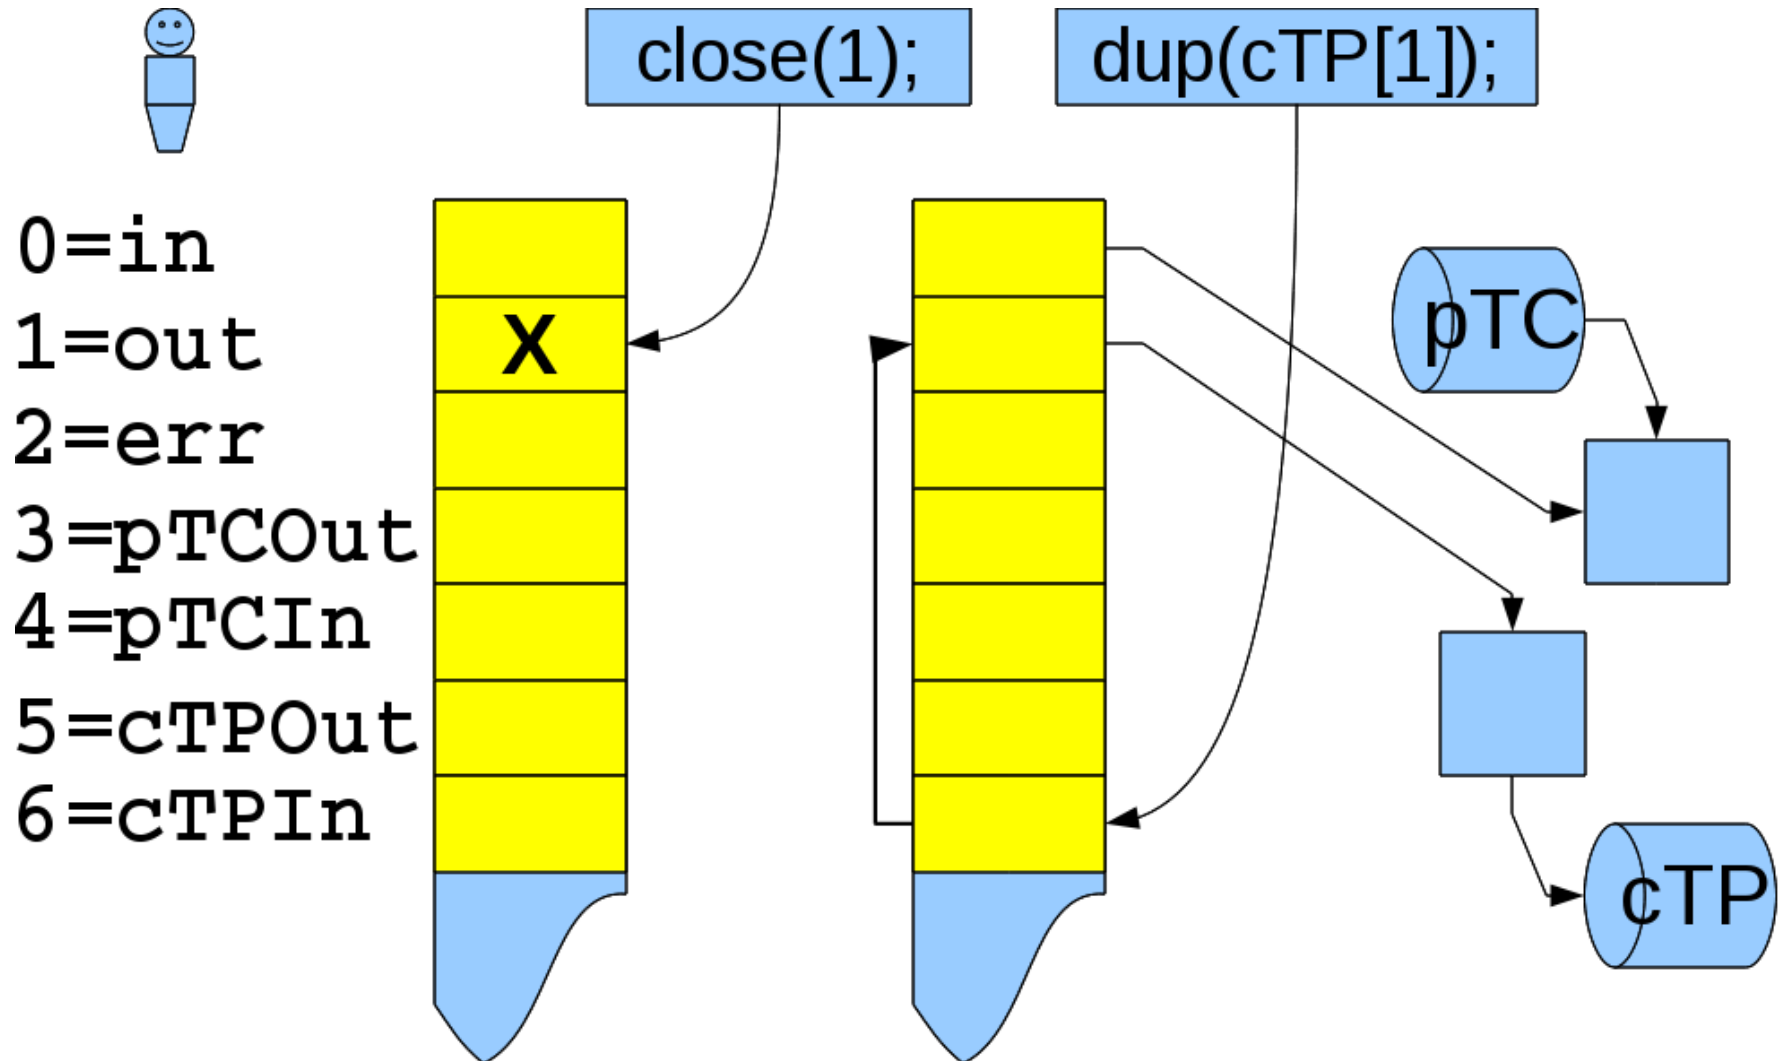
# dup( ) and pipe( ), 1



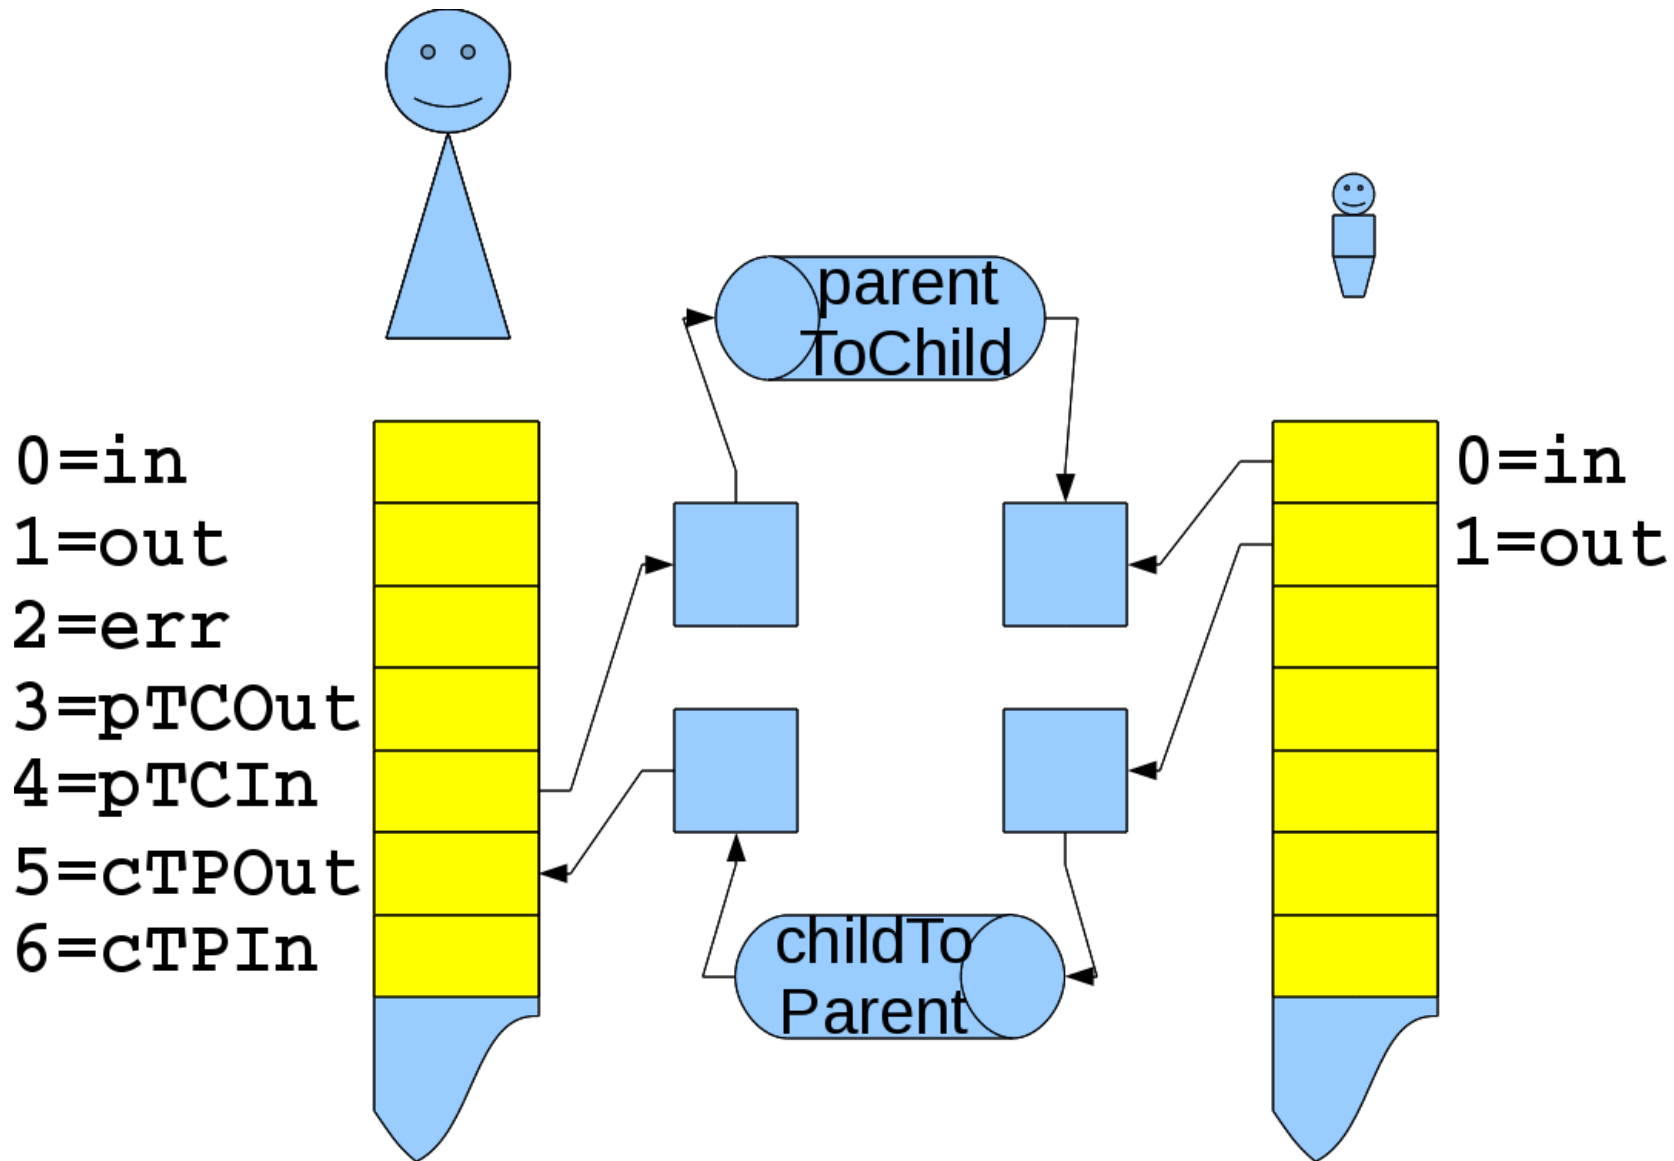
# dup( ) and pipe( ), 2



# dup( ) and pipe( ), 3



# dup( ) and pipe( ), 4



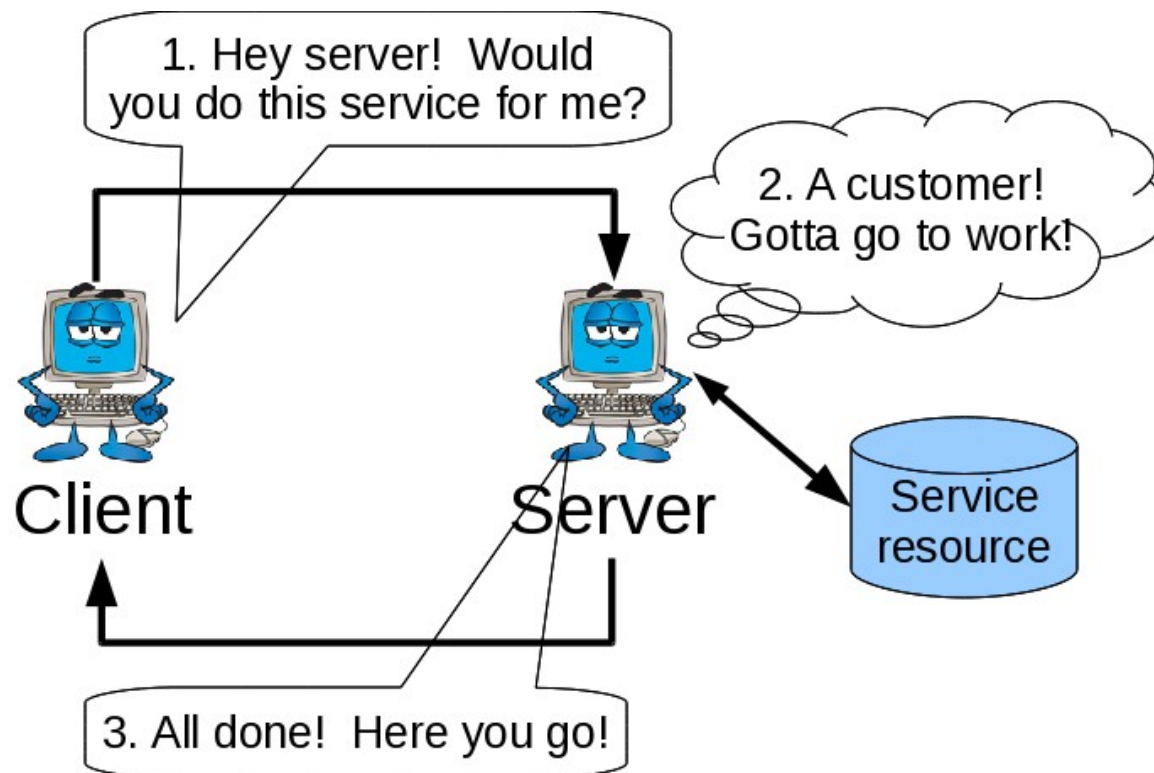
# Sockets

- ◆ Hey! Moving bytes to/from a file descriptor is such a ***grand idea*** that we can use it to move bytes to/from ***another process***
- ◆ Further, the process could be ***here*** (on the same machine)
- ◆ .....  
..... or way ***over here***  
(another machine across the Internet)
- ◆ We call this ***sockets!***

# The client-server computing model

1. Client asks server for a service
2. Server does service
3. Server returns result to client

Examples: `ssh`, `sftp`, `http`, *etc.*





# Processes talking to each other on different computers

Identify service by *IP address* and *port*

IP address: *Which computer?*

Humans like strings: “*www.depaul.edu*”

Computers like numbers: *75.102.246.202*

DNS: Domain Name Service

Given name get number (or vice versa)

Computers refer to themselves by the “loopback address”

*127.0.0.1* (integers)

*localhost* or *localhost.localdomain* (string)

# Ports:

Ports: *Which service on the given computer?*

Can range from 0 . . . 65535?

Common ones:

20 (ftp data), 21 (ftp control)

22 (ssh), 23 (telnet <-- DO NOT USE TO LOG IN!)

37 (time)

80 (www/http)

# Packets (or datagrams)

- Any communication is split up into manageable chunks (“packets”) that are sent individually.
- These chunks get routing, checksum, cryptographic, *etc.* info added to them

Packet from layer N+1



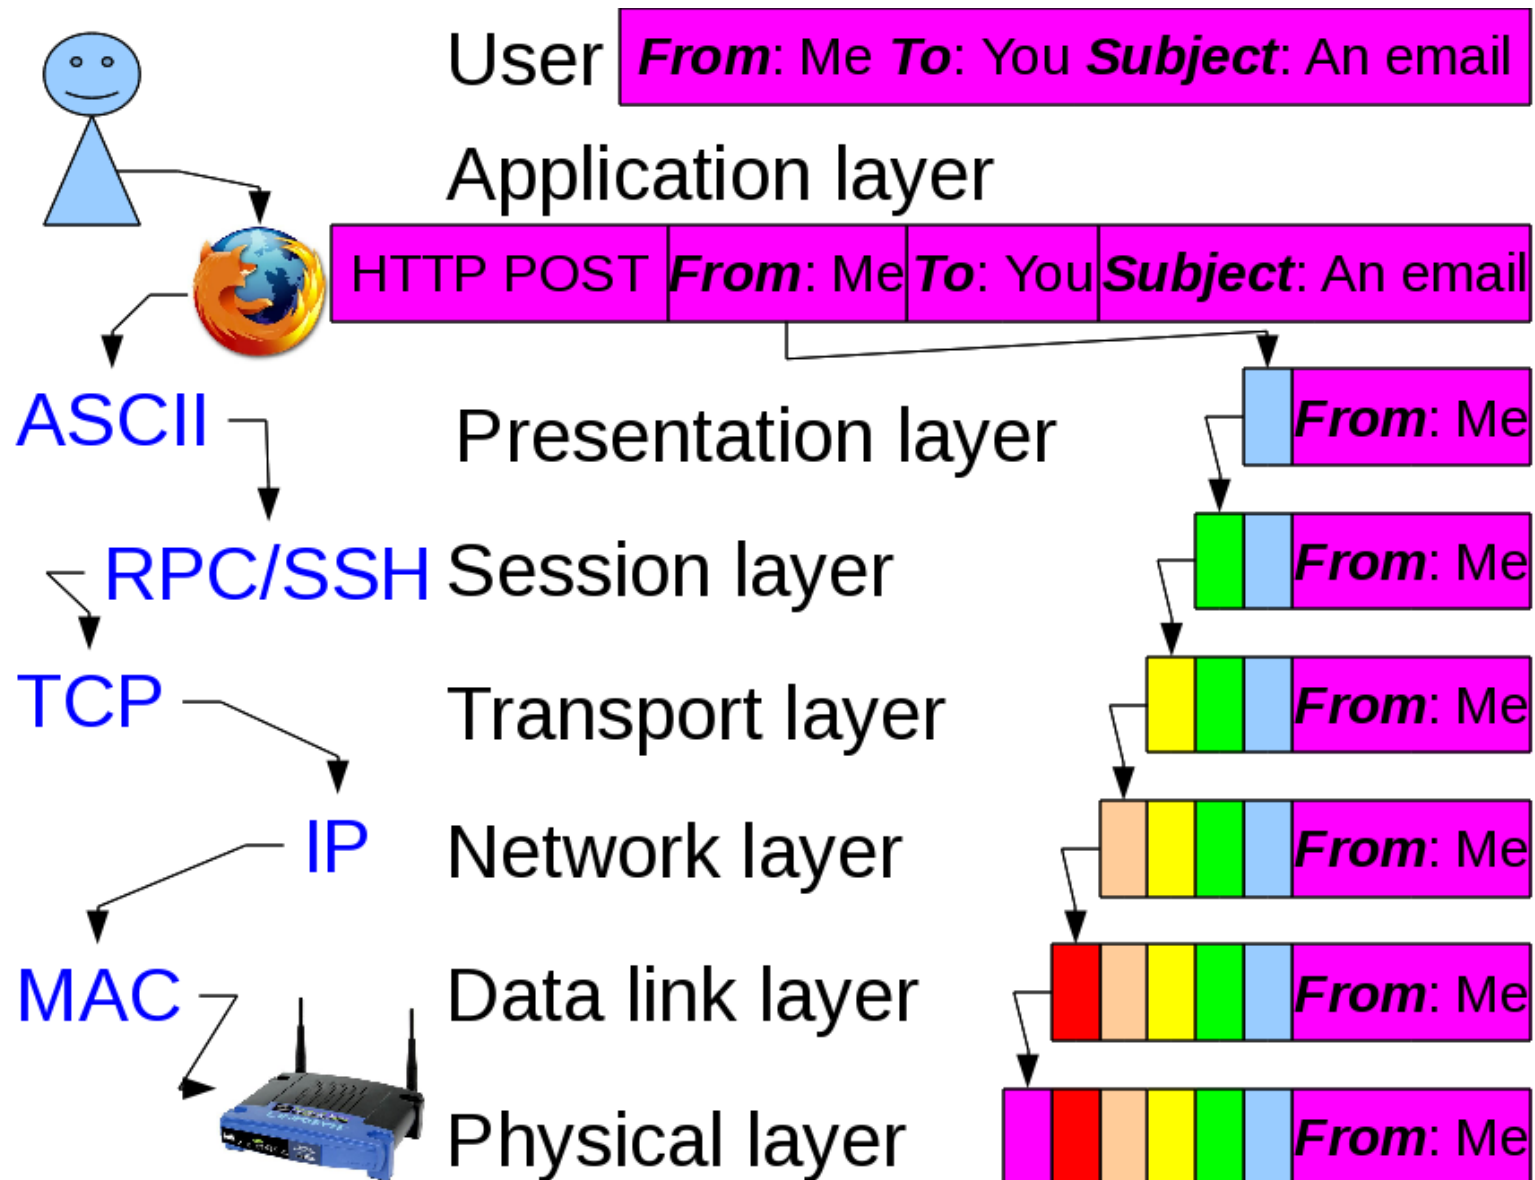
**Layer N** *“Let’s compute the cryptographic hash and append it”*



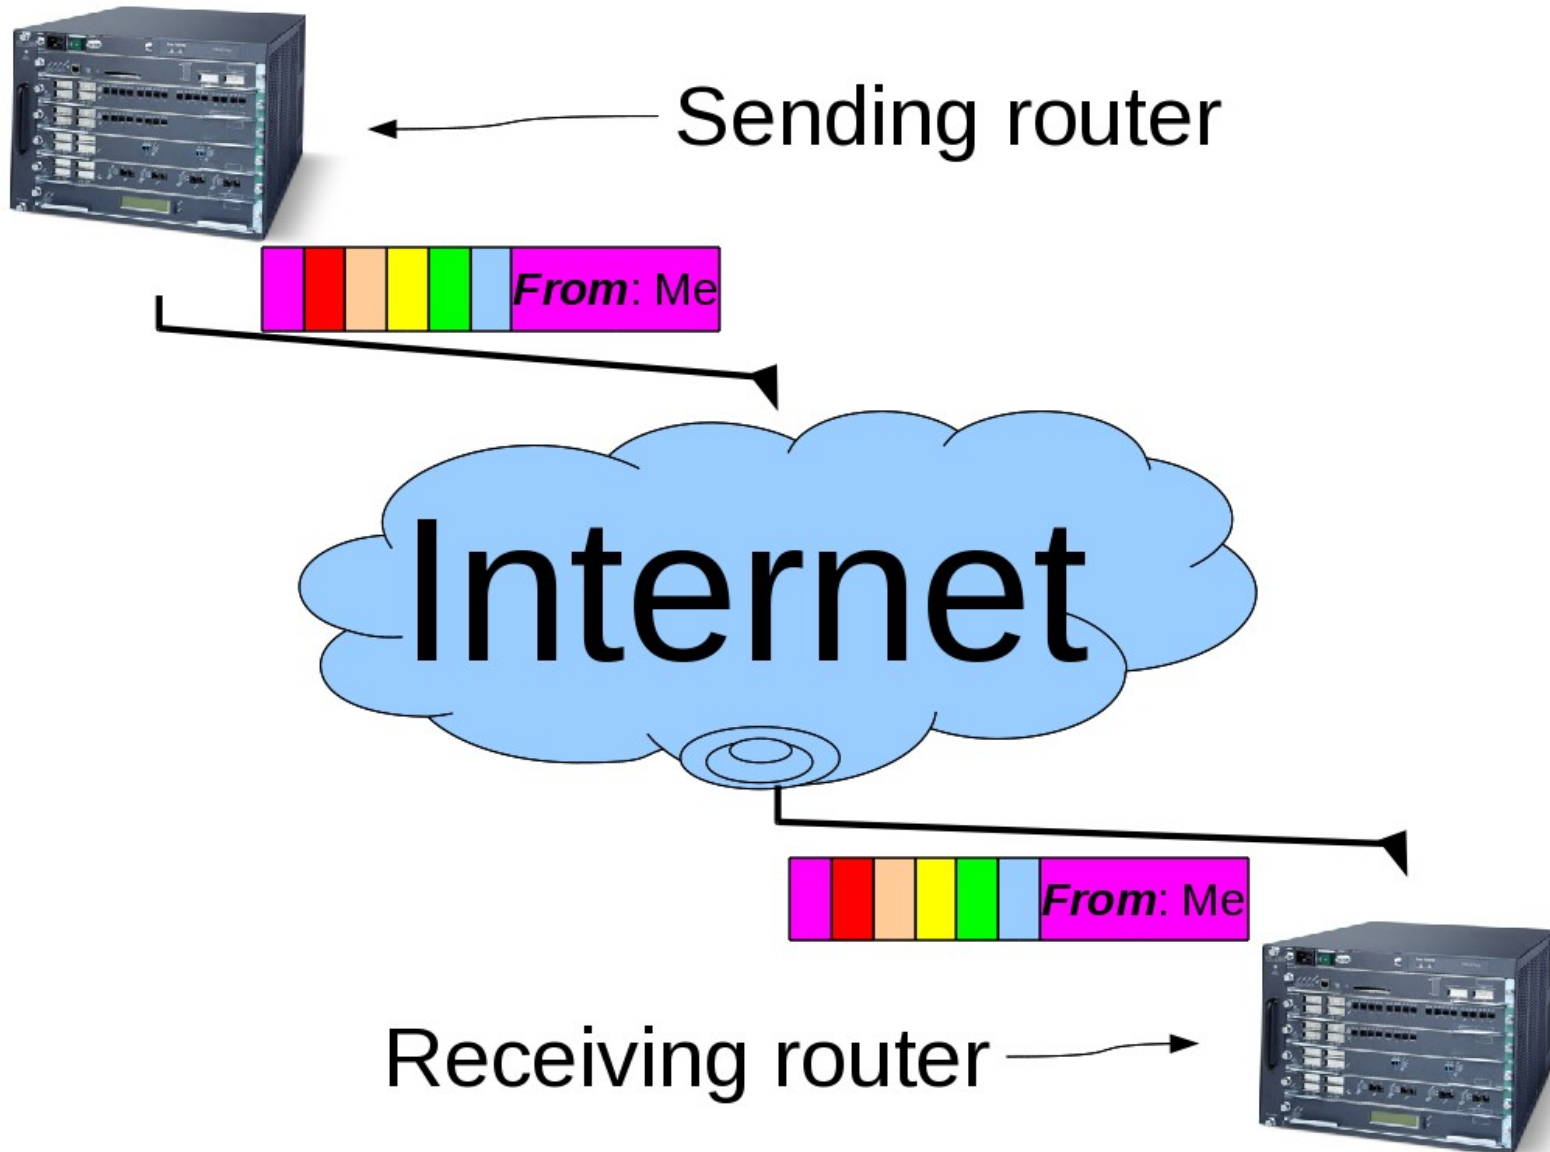
Layer N-1 appends its own header



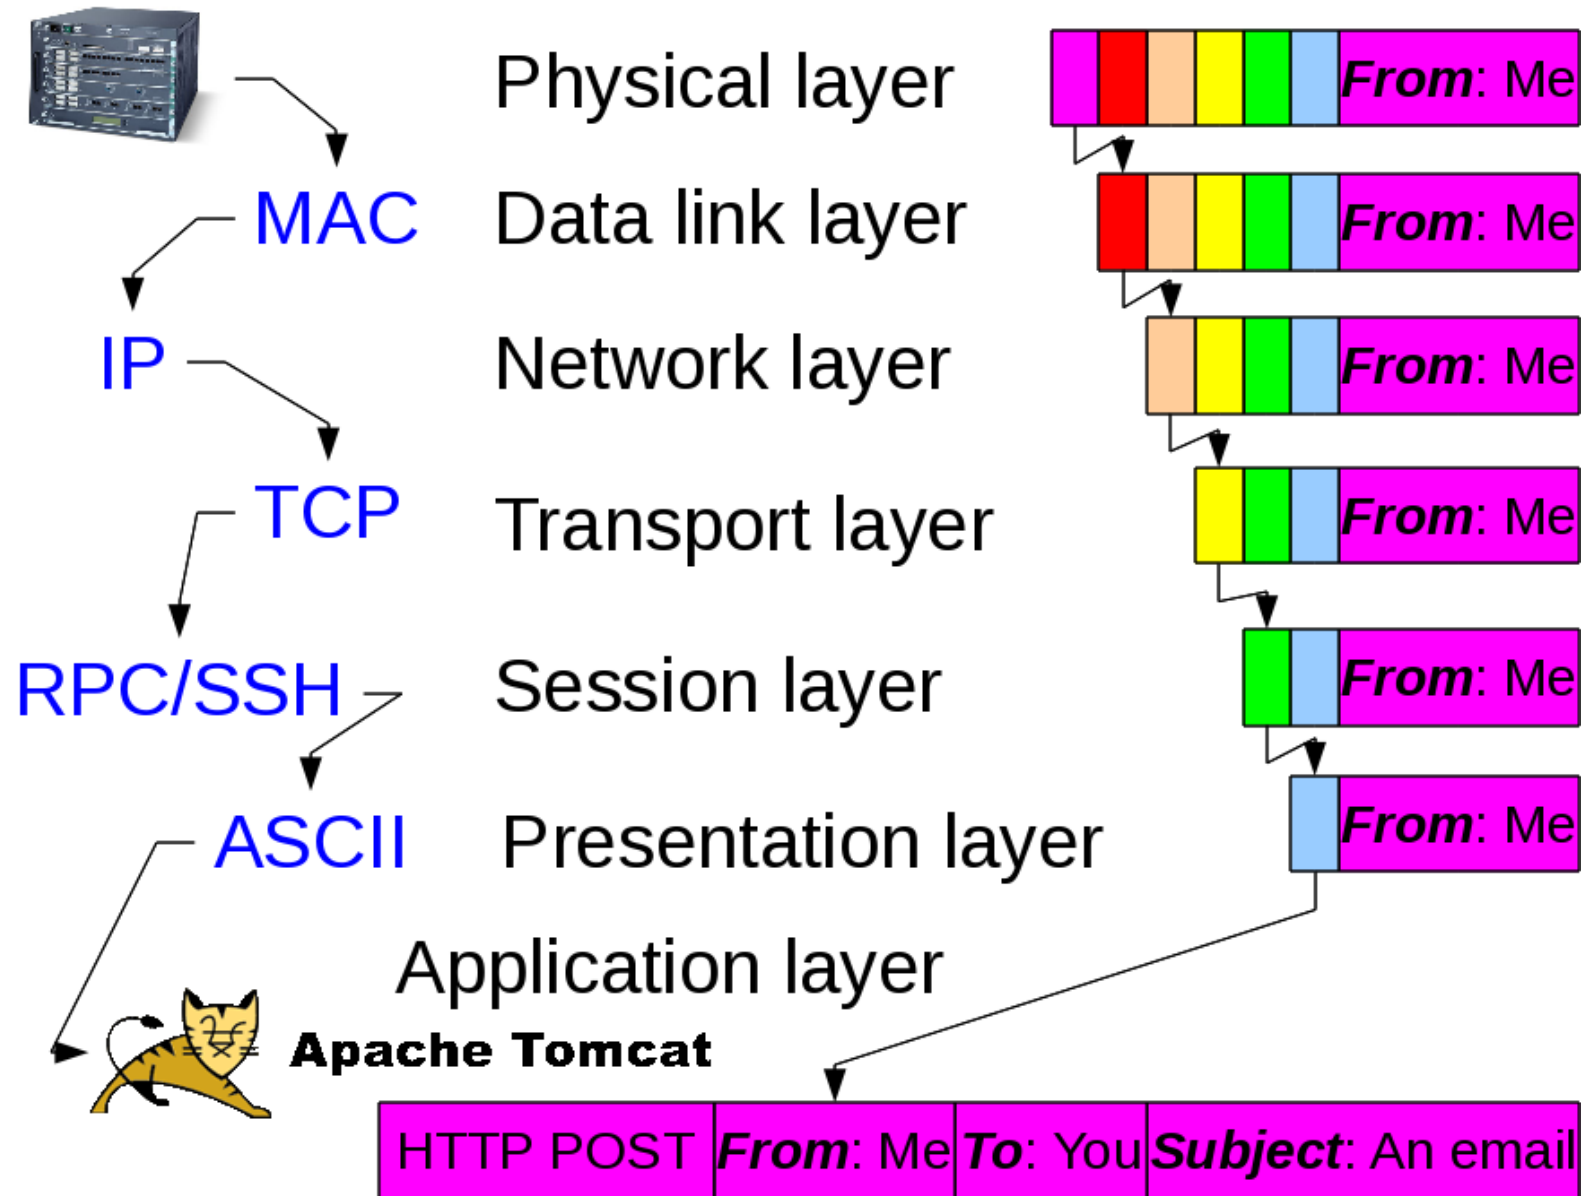
# So I know with whom I wanna talk, how do I communicate? (1)



# So I know with whom I wanna talk, how do I communicate? (2)



# So I know with whom I wanna talk, how do I communicate? (3)



# Tell me more about packets being sent over the 'net!

Two common networking protocols:

UDP: User Datagram Protocol

Fast! Unreliable!

*“Just send the packets! Missed the last one? Don't worry, here's another!”*

**Question: When would you use this?**

TCP: Transmission Control Protocol

Slower! Transmission is verified!

*“Just a received a packet. Slow down! Let's make sure it hasn't been corrupted and that I have them all.”*

**Question: When would you use this?**

# Sockets!

**Question:** Cool deal! How do I program it?

**Answer:** With sockets, silly!

## Socket communication

- Done with `read( )` and `write( )` (just like for files)
- Both have their own <address,port> pair
- Socket provides non-transient 2-way communication link



# Server Side:

```
#include <sys/socket.h> //For socket()  
#include <netinet/in.h> //For sockaddr_in and htons()  
#include <netdb.h>      //For getaddrinfo()  
#include <errno.h>       //For errno var  
#include <sys/stat.h>    //For open(), read(), write()  
#include <fcntl.h>       // and close()
```

- `socket()`: Ask OS for a socket
- `bind()`: Bind socket and port together
- `listen()`: Tell how many clients may queue
- `accept()`: Wait until a client connects
- `write()`: Write to client/server
- `read()`: Read from client/server
- `close()`: Close socket with client/server.

# Client Side:

```
#include <sys/socket.h> //For socket()  
#include <netinet/in.h> //For sockaddr_in and htons()  
#include <netdb.h>      //For getaddrinfo()  
#include <errno.h>      //For errno var  
#include <sys/stat.h>   //For open(), read(), write()  
#include <fcntl.h>      // and close()
```

- `getaddrinfo()`: Find server's IP address
- `socket()`: Ask OS for a socket
- `connect()`: Attempt to connect to server
- `write()`: Write to server
- `read()`: Read from server
- `close()`: Close socket with server.

# socket()

```
// Create a socket
int socketDescriptor =
    socket(AF_INET,          // AF_INET domain
          SOCK_STREAM,      // Reliable TCP
          0);
```

- Returns
  - A file descriptor that the server uses to see if a client has connected, or,
  - -1 on error
- There's also `SOCK_DGRAM` for UDP
- Last parameter type if used for `SOCK_RAW`

# bind()

// Bind socket to port

// We'll fill in this datastruct

sockaddr\_in socketInfo;

// Fill socketInfo with 0's

memset(&socketInfo, '\0', sizeof(socketInfo));

// Allow connections from myself only:

socketInfo.sin\_addr.s\_addr = inet\_addr("127.0.0.1");

// Allow machine to connect to this service

socketInfo.sin\_addr.s\_addr = INADDR\_ANY;

// Use std TCP/IP

socketInfo.sin\_family = AF\_INET;

// Tell port in network endian with htons()

socketInfo.sin\_port = htons(portNumber);

// Try to bind socket with port and other specifications

int status = bind(socketDescriptor, // from socket()

(sockaddr\*)&socketInfo,

sizeof(socketInfo)

);

status == -1 on error

# listen()

```
// Tell OS how many clients may queue  
up for this server  
int status =  
    listen(socketDescriptor,  
           maxNumPendingClients);
```

- (Almost) ready to listen to port!
- 5 is a good default for `maxNumPendingClients`.
- If `status == -1` then error

# accept()

```
// Accept connection to client  
int clientDescriptor =  
    accept(socketDescriptor, NULL, NULL);
```

- Wait (by default) for someone to actual connect
- Returns
  - a file descriptor for talking with one particular client, or
  - -1 for error
- `connectionDescriptor` for talking with that one client (there may be others for other clients)
- `socketDescriptor` is for listening to socket.

# Your turn!

**Question:** Hey! How is the server supposed to do two (or more!) things at once?

How do we get the server to both:

1. wait for another client to connect by listening to `socketDescriptor`, and
2. handle the current client(s) request by talking on `clientDescriptor`?

# Do you speak *BIG* or *little* Endian?

Now that we're talking . . . we'd better use same endian!

```
// Host to network long (ie. 32-bit)  
uint32_t htonl(uint32_t hostlong);
```

```
// Host to network short (ie. 16-bit)  
uint16_t htons(uint16_t hostshort);
```

```
// Network to host long (ie. 32-bit)  
uint32_t ntohl(uint32_t netlong);
```

```
// Network to host short (ie. 16-bit)  
uint16_t ntohs(uint16_t netshort);
```



# Your turn again!

- ◆ Write a server program that
  1. waits for a client to connect
  2. for any connected client it `read( )`s characters and `write( )`s the `toupper( )` of them.

Client-side time!

# getaddrinfo( )

```
// Get info on server (the "hostName")
int getaddrinfo
(const char*          hostName,
 const char*          service, //e.g. "ftp"
 const struct addrinfo *hints,
 struct addrinfo**    resultPtr);
// Also: getnameinfo( )
    Gets info on host given integers
```

- Sets `resultPtr` to datastructure with info on host `hostName`
  - `hostname` is C string, e.g. `"localhost.localdomain"`
  - returns integer: 0 == success, 0 != error.

# getaddrinfo( ) (derived from wikipedia example)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <netdb.h>
#include <netinet/in.h>
#include <sys/socket.h>

#ifdef NI_MAXHOST
#define NI_MAXHOST 1025
#endif

#define NAME_LEN 64
```

# getaddrinfo( ) (derived from wikipedia example)

```
int main(void)
{
    struct addrinfo *hostPtr;
    struct addrinfo *res;

    /* resolve domain name into a list of addrs */
    int haveValidName = 0;
    int error;

    while (!haveValidName)
    {
        char          name[NAME_LEN];

        printf("URL? ");
        fgets(name,NAME_LEN,stdin);
        char*         cPtr = strchr(name,'\n');
        if (cPtr != NULL)
            *cPtr = '\0';
```

# getaddrinfo( ) (derived from wikipedia example)

```
error = getaddrinfo(name, NULL, NULL, &hostPtr);
```

```
switch (error)
```

```
{
```

```
case 0 :
```

```
    haveValidName = 1;
```

```
    break;
```

```
case EAI_SYSTEM :
```

```
    perror("getaddrinfo");
```

```
    break;
```

```
default :
```

```
    fprintf(stderr, "getaddrinfo: %s\n",
```

```
            gai_strerror(error));
```

```
}
```

```
}
```

# getaddrinfo( ) (derived from wikipedia example)

```
/* loop over all returned results and do inverse lookup */
for (res = hostPtr; res != NULL; res = res->ai_next)
{
    char hostname[NI_MAXHOST] = "";

    error = getnameinfo(res->ai_addr, res->ai_addrlen, hostname,
NI_MAXHOST, NULL, 0, 0);
    if (error != 0)
    {
        fprintf(stderr, "error in getnameinfo: %s\n", gai_strerror(error));
        continue;
    }
    if (*hostname != '\0')
        printf("hostname: %s\n", hostname);
}

freeaddrinfo(hostPtr);
return 0;
}
```

# connect()

```
// Connect to server
sockaddr_in server;

// Clear server datastruct
memset(&server, 0, sizeof(server));
// Use TCP/IP
server.sin_family      = AF_INET;
// Copy connectivity info from info on server ("hostPtr")
server.sin_addr.s_addr =
    ((struct sockaddr_in*)hostPtr->ai_addr)->sin_addr.s_addr;
// Tell port # in proper network byte order
server.sin_port        = htons(portNumber);

int status = connect(socketDescriptor,&server,sizeof(server));
```

◆ -1 means error



# read(), write() and close()

As previously stated:

```
// Read from file/socket
// numRead==0 means "EndOfFile", numRead==-1 means "error"
int numRead =
    read(connectDescriptor,bufferAddress,bufferLen);
int numRead =
    recv(connectDescriptor,bufferAddress,bufferLen,int flags);

// Write to file/socket: numWritten == -1 means "error"
int numWritten =
    write(connectDescriptor,bufferAddress,bufferLen);
int numWritten =
    send(connectDescriptor,bufferAddress,bufferLen,int flags);

// Close connection: status == -1 means "error"
int status = close(descriptor);
```

# But sometimes you don't want to wait for socket input

```
int recv(int connectDescriptor, void*  
    bufferPtr, int bufferLen, int flags)
```

Reads up to `bufferLen` bytes into the buffer pointed to by `bufferPtr` from file descriptor `connectDescriptor`. `flags` tells how to read, where `MSG_DONTWAIT` means "non-blocking".

Returns number of bytes read, or returns `-1` and sets global var `errno` to `EAGAIN` if the flag was `MSG_DONTWAIT` and there was nothing to read.

# Short counts with `recv( )`

- Short counts occur during:
  - Encounter End of file (EOF) when reading file (expected)
  - Reading text from terminal (also expected)
  - Reading from network or pipes if get interrupted by catching any sort of signal (an annoyance!)
  - **Question:** Did the fact that `read( )` or `recv( )` returned something mean that it ***got*** something, or that it ***was interrupted?***
- ***Oh no! Can nothing save us?!?!?***

# Robust I/O Package to the rescue!

```
/* From authors' thread-safe, buffered I/O package.
Same interface as read() */
ssize_t rio_readn(int fd, void *usrbuf, size_t n)
{
    size_t  nleft = n;
    ssize_t nread;
    char*    bufp = usrbuf;

    while (nleft > 0) {
        if ((nread = read(fd, bufp, nleft)) < 0) {
            if (errno==EINTR) /* Interrupted by sig handlr rtn? */
                nread = 0;    /* Yes: Call read() again */
            else
                return -1;    /* No:  Have some other error */
        }
        else if (nread == 0) /* Have EOF? */
            break;           /* Yes: Just quit loop */
        nleft -= nread;      /* Else that many fewer chars to get */
        bufp += nread;       /* Advance in buffer to read more */
    }
    return (n - nleft);      /* For non errors return val >= 0 */
}
```

# Your turn!

- ◆ Write a client program that:
  1. Connects with the server
  2. Asks the user for text
  3. Sends the text to the server
  4. Gets the response back from the server and prints it

**Next time**

ncurses cursor control